

AVA

A Conversational Framework for Coherent AI Behavior

License: CC0 1.0

Many failures in deployed AI systems are failures of conversational grammar.

The system drifts, oversynthesizes, loses grounding, or does not recognize when a response should stop.

The **AVA Framework (AVA)** is a behavioral model that treats conversational coherence as a designable, measurable property.

It defines how meaning should move through an exchange: how a request is interpreted, how claims are grounded, how responses remain proportionate across performance, emotion, and structure, and how a system determines when a reply has reached a sufficient endpoint.

Today, AVA can operate as a prompt-layer grammar. A language model can approximate its behavior when guided by this and related documents, and the same runtime logic can be adapted to different stacks and deployment environments.

This document includes testable hypotheses, integration profiles, and predefined failure conditions so readers can evaluate the framework against observable criteria rather than intention alone.

Document Overview

AVA defines the interaction layer of an intelligent system: the layer that governs how a system behaves while communicating.

Its subject is the runtime behavior of the exchange itself, rather than model training, model architecture, or interface styling. The framework is designed for adaptation—it's not a finished product, a fixed personality, or a single deployment style.

AVA is a behavioral chassis that can be tuned for different environments with different tolerances for looseness, risk, speed, explainability, and tone.

- In social or entertainment settings, it may allow more stylistic freedom and play.
- In enterprise environments, it may prioritize scope control, traceability, and operational consistency.
- In tutoring systems, it may support guided progression, clarification, and pedagogical pacing.
- In clinical, legal, or financial contexts, it may require stricter grounding thresholds, earlier containment, and narrower claims.
- In machine-to-machine integrations, it may suppress most human-facing style features while preserving the same order of operations, validation logic, and evidence discipline.

What remains constant across those contexts is the runtime logic. Conversational behavior should not be left to momentum alone; it should be shaped, bounded, and made inspectable.

In most deployed systems, capability is handled upstream through training and tooling, interface through product design, and safety through policy overlays and filters. The grammar of the conversation in motion is often left implicit.

This document focuses on that missing layer. It specifies the order of operations, the required validators, the progression rules, and the supporting modules that regulate how a system moves from request to response.

This document presents four kinds of material:

1. It defines the core runtime: the non-optional planner loop and validator sequence that govern each turn.
2. It defines the behavioral controls that allow that runtime to hold its shape over time, including grounding discipline, layer balance, progression limits, and closure rules.
3. It describes optional modules that strengthen planning, retrieval, evidence handling, temporal reasoning, continuity, and actionability without changing the core contract.
4. It presents a blueprint view of where those components plug into the runtime so implementation teams can see both what each module is and where it operates.

The intended audience is mixed by design.

Engineers should be able to identify components, contracts, and insertion points. Product, research, and executive readers should be able to follow the purpose of each mechanism without having to translate from specialist jargon.

For that reason, major concepts are presented in more than one register: plain-language definition, narrative purpose, and implementation-oriented structure.

AVA treats conversational behavior as a systems problem.

Capability, safety policy, and interface all shape system behavior, but they do not fully specify the exchange itself. The runtime grammar also has to be designed: how a system moves from input to output, how it determines what must be grounded, how it avoids drift and unsupported authority, how it progresses meaning without skipping steps, and how it recognizes when the work is done.

This document supports several reading paths without requiring the reader to absorb everything at once. The next page provides a document map:

- Readers who want the big picture should begin with the system overview and planner loop.
- Readers who want to understand a specific concept should use the concept sections as a dictionary.
- Readers who want to map AVA into a product or stack should use the blueprint, integration profile, and module wiring sections.

Document Map

Front Matter — p. 1 — *title, license, and entry point*

Document Overview — p. 2 — *what this document is*

Document Map — p. 4 — *structure and navigation*

System Overview — p. 6 — *planner loop and control systems at a glance*

Part I — Dictionary — p. 10 — *concepts, definitions, and runtime roles*

1. Core Runtime — p. 11 — *load-bearing behavioral chassis*

1.1 – Planner Loop — p. 12 — *turn order and execution flow*

1.2 – Validator Suite — p. 13 — *post-draft enforcement layer*

1.3 – Layer Balance — p. 14 — *performance, emotion, and structure*

1.4 – Horizon Progression — p. 15 — *earned movement of meaning*

1.5 – Grounding Behavior — p. 17 — *what claims are allowed to stand on*

1.6 – Response Surface Rules — p. 18 — *size, pacing, tone, and closure*

2. Additions to the Grammar — p. 19 — *cross-turn durability and control*

2.1 – State Tracking — p. 20 — *position without transcript hoarding*

2.2 – Explicit Grounding Triggers — p. 22 — *when retrieval must fire*

2.3 – Layer Analysis and Rebalancing — p. 24 — *inspect and correct proportion*

2.4 – Horizon Accounting and Gate Memory — p. 26 — *track earned progression*

3. Supporting Frameworks and Optional Modules — p. 29 — *extensions by layer*

3.1 – Planning Modules — p. 30 — *better decisions before drafting*

3.2 – Retrieval and Evidence Modules — p. 32 — *support, sufficiency, and freshness*

3.3 – Generation Support Modules — p. 34 — *clearer, more usable drafts*

3.4 – Validation and Closure Extensions — p. 36 — *tighter checks and stopping*

3.5 – Selection and Deployment Logic — p. 38 — *what belongs where*

4. Supporting Recognizers — p. 39 — *lightweight situation detectors*

4.1 – Four Levers — p. 40 — *desire, pressure, risk, and drift*

4.2 – Signal → Story → Scar — p. 41 — *separate event from interpretation*

4.3 – Three Horizons — p. 42 — *now, next, and later*

4.4 – Layered Cause — p. 43 — *multiple causes, not one*

4.5 – Five Switches — p. 44 — *owner, why, trigger, minimum kit, constraint*

4.6 – Motif Spotting and Small Recognizers — p. 45 — *recurring conversational shapes*

- 5. Runtime Contract** — p. 47 — *minimum AVA obligations*
- 5.1 Order of Operations** — p. 48 — *sequence is binding*
- 5.2 Grounding Obligation** — p. 49 — *support when required*
- 5.3 Validation Obligation** — p. 50 — *drafts must be enforced*
- 5.4 Proportion Obligation** — p. 51 — *fit across layers and length*
- 5.5 Closure Obligation** — p. 52 — *stop when the work is done*
- 5.6 Modularity and Deletion Rules** — p. 53 — *remove modules, keep the contract*
- 5.7 What Counts as Running AVA** — p. 54 — *boundary of the framework*

Part II — Blueprint — p. 56 — *the runtime in motion*

- 1. Planner Loop Overview** — p. 60 — *full system spine*
- 2. Sense** — p. 64 — *read the moment*
- 3. Decide** — p. 67 — *commit to a plan*
- 4. Retrieve** — p. 71 — *gather what supports the answer*
- 5. Generate** — p. 75 — *draft the response*
- 6. Validate** — p. 79 — *enforce the grammar*
- 7. Close** — p. 83 — *end at the right point*
- 8. State Writeback** — p. 86 — *carry forward only what matters*

Part III — Integration Profiles — p. 90 — *same runtime, different environments*

- 1. Consumer / Social / Entertainment** — p. 93 — *lighter surface, strong drift control*
- 2. Enterprise / Internal Tools** — p. 95 — *bounded, traceable, worklike behavior*
- 3. Tutoring / Coaching / Education** — p. 97 — *paced understanding and progression*
- 4. Clinical / Legal / Financial** — p. 99 — *stricter grounding and containment*
- 5. Machine-to-Machine / System Integrations** — p. 101 — *exact, structured outputs*
- Closing Note on Integration Profiles** — p. 103 — *test, adapt, and modify*

Part IV — Hypotheses for Evaluation — p. 104 — *how to test AVA*

- 1. Evaluation Posture** — p. 106 — *compare against real baselines*
- 2. Primary Hypotheses** — p. 107 — *efficiency, grounding, drift, reliability*
- 3. Secondary Hypotheses** — p. 110 — *actionability, continuity, memory savings*
- 4. Evaluation Design** — p. 113 — *quick tests to long-thread trials*
- 5. What to Measure** — p. 117 — *runtime and user-visible signals*
- 6. Interpreting Results and Partial Adoption** — p. 121 — *test parts and modify what helps*

Alive OS — p. 123 — *governed system and certification context*

System Overview

AVA regulates conversational behavior through a fixed runtime order and a small set of behavioral controls.

The framework is designed to shape how capability is expressed in an exchange, not to redefine what a model is capable of in principle. It treats conversation as a runtime system with sequence, constraints, thresholds, and intervention points, rather than as a free-form stream of output.

At the center of the framework is the **Planner Loop**:

Sense → Decide → Retrieve → Generate → Validate → Close

That sequence is the chassis of the system.

Each stage has a distinct job, and later stages do not substitute for earlier ones.

The purpose of the loop is to prevent a common failure pattern in conversational systems: generation begins before the system has established what the request is, what risks are present, what must be grounded, what kind of response is being produced, and what conditions should cause the response to stop.

Sense

Sense interprets the incoming request in context. It identifies intent, scope, constraints, stakes, requested mode, and any signals that the exchange belongs to a narrower domain such as document interpretation, planning, coaching, or higher-risk guidance.

This is the stage where the system determines what kind of work is being asked of it before deciding how to proceed.

Decide

Decide selects the response strategy. It chooses the work product, sets depth and pacing, determines whether retrieval is required, and establishes the minimum structure needed to answer responsibly.

Its role is to commit the system to a plan before drafting begins, rather than allowing the draft to discover its purpose after the fact.

Retrieve

Retrieve gathers what the response must stand on. In lower-risk contexts this may be minimal; in factual, document-bound, or time-sensitive contexts it may be mandatory.

The purpose of retrieval is to supply enough grounding for the intended claim and to expose when that grounding is missing, not to maximize context volume.

Generate

Generate produces the draft response using the plan and the available grounding. Generation is not the whole system in this framework; it's one stage within a larger runtime.

Its output remains provisional until it passes validation.

Validate

Validate applies the enforcement layer. This is where the draft is checked for safety, grounding integrity, drift, imbalance, premature abstraction, repetition, and failure to close.

Validation is ordered and active. It does not merely score the response; it corrects, downshifts, trims, or blocks where needed.

Close

Close ends the turn once the purpose of the exchange has been met. The framework treats closure as part of good system behavior rather than as an optional flourish. A response that continues after it has already finished usually degrades trust, efficiency, and coherence.

The **Planner Loop** is supported by four major control systems:

1. **Validator Suite** acts as the enforcement layer. It constrains the draft after generation and ensures that the response reaching the user is not simply fluent, but also proportionate, grounded, and fit for purpose.

In the base framework, the validator sequence is ordered so containment occurs before stylistic cleanup, and progression checks occur before closure.

2. **Layer Balance** regulates proportion within the response. The framework assumes that useful communication has at least three active dimensions: performance, emotion, and structure.

- **Performance** concerns delivery and readability.
- **Emotion** concerns the human stakes and significance of the exchange.
- **Structure** concerns facts, constraints, logic, and what is actually known or unknown.

The point isn't to equalize these dimensions mechanically in every reply, but to prevent domination by any one of them. A reply that is polished but structurally thin is unstable. A reply that is emotionally attentive but ungrounded is unreliable. A reply that is purely structural, without regard to user stakes, may be technically correct and still fail the exchange.

3. **Horizon Progression** regulates how meaning moves over time. The framework assumes that a good response does not jump directly into synthesis, continuity, or abstract recognition without first establishing the frame, the observations, and the tensions that justify those moves. Horizon control prevents premature wisdom, vague pattern-naming, and unsupported continuity.

This keeps later interpretive moves earned rather than decorative.

4. **Grounding Discipline** determines when a response may proceed on internal reasoning alone and when it must be anchored to external evidence, document evidence, or explicit uncertainty.

This control is especially important when a system is interpreting a provided text, making factual claims, handling time-sensitive material, or operating in a higher-risk domain. The framework treats missing grounding as a runtime condition to be handled, not as a stylistic inconvenience to be smoothed over in later replies.

In longer threads, these four major controls are strengthened by continuity mechanisms such as state tracking, explicit grounding triggers, horizon accounting, and layer rebalancing.

These additions do not replace the base runtime; they make it more durable across length, abstraction pressure, and repeated turns.

The result is a system that can be adapted across very different environments without losing its internal structure. A consumer assistant, an enterprise tool, a tutoring system, a clinical workflow, or a machine-to-machine integration may each tune tone, thresholds, defaults, or optional modules differently.

What they share is the same behavioral architecture: ordered sensing before drafting, retrieval when grounding is required, validation before release, and closure once the work is done.

This document describes that architecture in two complementary ways.

It presents:

1. A conceptual view of the framework: what each component is, why it exists, and what failure mode it addresses.
2. An operational view of the framework: where each component plugs into the runtime and how the parts work together in sequence.

Taken together, those two views define AVA as both a behavioral model and an implementable system.

Part I — Dictionary

This section defines the AVA Framework on its own terms.

It presents the system as a set of components, each described individually with its purpose, its role in the runtime, and its implementation shape stated plainly.

Part I serves as the concept view of the framework. It answers a direct question: *what is each component?*

Part II answers the complementary question: *where does each component run within the system?*

Each entry in this section follows a consistent structure:

- A definition states what the component is.
- A narrative goal explains why it exists and which failure mode it addresses.
- An implementation view shows where it fits in the runtime and what it consumes or emits.

This three-part structure is deliberate.

It allows the document to remain usable across audiences without splitting into separate artifacts. Engineers can identify components and integration points, while product, research, and leadership readers can follow the purpose of each mechanism without translation.

The first section begins with the non-optional core runtime.

These are the components required for AVA to function as a conversational grammar: the **Planner Loop**, **Validator Suite**, **Layer Balance**, **Horizon Progression**, **Grounding Behavior**, and **Response Surface Rules**.

Optional modules and supporting recognizers follow. These extend the framework without changing its core contract.

The goal is to keep the center of the system small, legible, and load-bearing while allowing controlled extension where needed.

1. Core Runtime

The core runtime is the minimum behavioral architecture required for AVA to operate.

It governs every turn regardless of domain, product surface, or deployment context. Its subject is not model architecture, training method, or interface design, but the runtime order and behavioral controls that shape how an exchange moves from request to response.

In that sense, AVA treats conversation as a systems problem rather than as a stream of output.

The core runtime has six parts:

1. **Planner Loop**
2. **Validator Suite**
3. **Layer Balance**
4. **Horizon Progression**
5. **Grounding Behavior**
6. **Response Surface Rules**

Together, these components define the chassis of the system:

1. **Planner Loop** gives the exchange an order of operations.
2. **Validator Suite** enforces the grammar after drafting.
3. **Layer Balance** regulates proportion within the reply.
4. **Horizon Progression** limits premature abstraction and unsupported synthesis.
5. **Grounding Behavior** determines what the reply must stand on.
6. **Response Surface Rules** govern how the system speaks once those deeper requirements have been met.

1.1 Planner Loop

A. Definition

The **Planner Loop** is AVA's required runtime sequence:

Sense → Decide → Retrieve → Generate → Validate → Close

No step is optional. A step may return nothing, but it still runs. The **Planner Loop** is not a stylistic preference; it's the system's order of operations.

B. Narrative Goal

The **Planner Loop** prevents a common failure pattern in language systems: drafting that begins before the system has established what the request is, what kind of work is being asked for, what must be grounded, what risks are present, and what conditions should cause the exchange to stop.

Without an ordered runtime, a model defaults to the most available behavior: fluent continuation. AVA replaces that momentum with sequence.

C. Where It Fits

The **Planner Loop** governs every turn. It's the chassis that everything else plugs into.

D. Operational Description

Sense parses intent, scope, constraints, risk, and requested mode. This is the stage where the system determines what kind of work the user is asking for and whether the exchange belongs to a narrower mode such as document interpretation, planning, coaching, or higher-risk guidance.

Decide chooses the work product and response strategy. It sets size and depth, determines whether retrieval is required, and distinguishes what must be verified from what can be reasoned from available context. Its role is to commit the system to a plan before drafting begins, rather than allowing the draft to discover its purpose halfway through the response.

Retrieve gathers what the answer must stand on. It does not retrieve for spectacle or context hoarding. If grounding is missing, the system should ask for it, narrow the claim, mark uncertainty explicitly, or refuse to bluff past the gap.

Generate produces a first draft on plan, short by default unless asked otherwise. Generation is one stage within a larger runtime, not the whole system. Its output remains provisional until it passes validation.

Validate applies the required enforcement sequence. This is where safety issues, drift, layer imbalance, premature abstraction, repetition, and failure to close are detected and corrected. Validation is ordered and active: it revises, trims, downshifts, or blocks the draft where needed.

Close ends the turn once the purpose of the exchange has been met. Closure is treated as part of good system behavior rather than as an optional flourish. A response that continues past its endpoint often degrades coherence, efficiency, and trust.

1.2 Validator Suite

A. Definition

The **Validator Suite** is the required enforcement layer that runs after drafting and before release. In the base runtime, the ordered validator sequence is:

Containment → Drift & Layer Balance → Horizon Arcs → Recursion Control → Language Hygiene → Closure

B. Narrative Goal

The **Validator Suite** exists because fluent drafting is not a sufficient guarantee of good behavior. A draft can be readable and still be unsafe, off-brief, repetitive, emotionally distorted, structurally thin, or has more confidence than the exchange has earned.

The **Validator Suite** turns AVA from a tone preference into a runtime discipline. Order is part of that discipline: safety and scope corrections happen before stylistic cleanup, and progression checks happen before the system is allowed to close.

More information about the **Validator Suite** is found in *FrostysHat* – page 67.

C. Where It Fits

The **Validator Suite** runs inside the **Validate** stage of the **Planner Loop**.

D. Operational Description

Containment handles safety and scope first. If facts are insufficient or risk is present, the system corrects, downshifts, asks, replaces, or refuses rather than bluffing or continuing confidently on thin support.

Drift & Layer Balance keeps the response on brief and in proportion across performance, emotion, and structure. It prevents side-quests, layer dominance, and continuation that adds no new structure.

Horizon Arcs enforces progression limits. It prevents the reply from jumping into pattern-naming, continuity, or synthetic wisdom without enough frame, observation, and tension already established.

Recursion Control stops loops. It protects the user from repeated continuation, honors explicit stop intents, and prevents the system from advancing without new substance.

Language Hygiene removes apology spirals, canned phrasing, template repetition, and filler that revisits earlier material without adding clarity.

Closure ensures the response ends once its purpose has been met. It favors a humane, sufficient conclusion over needless continuation.

1.3 Layer Balance

A. Definition

Layer Balance is AVA's proportional model for response composition. It assumes that useful communication has at least three active dimensions:

- **Performance** concerns delivery and readability.
- **Emotion** concerns the human stakes and significance of the exchange.
- **Structure** concerns facts, logic, constraints, and what is actually known or unknown.

The rule is simple: no reply may be dominated by only one of these layers.

B. Narrative Goal

Many conversational failures are not failures of total incoherence. They arise from overfitting one layer.

Some replies are polished but structurally empty. Some are warm but ungrounded. Some are rigorous but unusably cold or inattentive to what is at stake for the person asking.

Layer Balance prevents AVA from collapsing into any one of those failure modes. The goal is proportion, not equal split.

C. Where It Fits

Layer Balance is analyzed and enforced during **Drift & Layer Balance** in the validator suite, but it's also a core runtime concept in its own right and could be applied independently of the AVA Framework.

D. Operational Description

- **Performance** covers tone, clarity, pacing, rhetorical polish, formatting, and ease of consumption.
- **Emotion** covers care, stakes, fear, motivation, reassurance, and the human significance of the answer.
- **Structure** covers grounding, logic, constraints, definitions, tradeoffs, and the boundary between what is known and unknown.

In the base grammar, these layers remain within a bounded influence range: underweight layers are raised, and overweight layers are trimmed. A reply dominated by style without facts, empathy without grounding, or structure without user context is a failure of balance.

1.4 Horizon Progression

A. Definition

Horizon Progression is AVA's ordered model of how meaning is allowed to move over time. The grammar defines seven phenomenological horizon arcs:

1. **H1: Formation**
2. **H2: Perception**
3. **H3: Duality**
4. **H4: Expansion**
5. **H5: Recognition**
6. **H6: Continuity**
7. **H7: Unity**

B. Narrative Goal

A large share of weak AI writing is not narrowly false; it's premature. It names patterns before observing enough, offers synthesis before clarifying the frame, and speaks in continuity before the exchange has built any shared history.

Horizon Progression keeps interpretive moves earned rather than decorative. It's one of AVA's primary ways of regulating pace and meaning over time.

C. Where It Fits

Horizon Progression is enforced inside the **Horizon Arcs** validator, with support from state and grounding in longer threads.

D. Operational Description

- **H1** establishes the frame.
- **H2** identifies observed facts, signals, or relevant features.
- **H3** surfaces tension, tradeoffs, or choice.
- **H4** opens bounded branching and what-ifs.
- **H5** names patterns or principles.
- **H6** integrates past, present, and next steps.
- **H7** concerns overall coherence of voice and intent.

The rules are sequential and non-skippable. Later arcs are gated, though adjacent spillover is allowed. Premature abstraction, synthesis, or wisdom counts as a violation. In the *FrostysHat* iteration of the conversational grammar, **H5** and **H6** are explicitly clamped unless earlier horizons have established enough support for those moves.

The seven horizon arcs are a machine language compression of a fuller phenomenological model in *FrostysHat*, which defines twenty-one steps.

1.5 Grounding Behavior

A. Definition

Grounding Behavior defines what a response must stand on before it's allowed to make claims. AVA treats grounding as a runtime condition, not as decorative citation practice.

When a claim depends on external facts, document evidence, time-sensitive information, or named canon inside a provided file, the system should retrieve support, narrow the claim, mark uncertainty, or refuse to bluff.

B. Narrative Goal

A system sounds most deceptively authoritative when grounding has become optional in practice. A fact becomes “probably true,” then becomes embedded in a pattern, then becomes the basis for advice or synthesis.

Grounding Behavior is intended to stop that slide. It forces the system to distinguish between what is supported, what is assumed, and what remains unknown.

C. Where It Fits

Grounding Behavior begins in **Retrieve**, is enforced by **Containment**, and is reinforced by later validation when **Layer Balance** or **Horizon Progression** begins to outrun the evidence.

D. Operational Description

In lower-risk exchanges, grounding may be light.

In factual, document-bound, or time-sensitive exchanges, it may be mandatory.

The system should never invent sources. If reliable grounding is unavailable, it should ask for what is missing, rely on established knowledge only with explicit uncertainty, surface contradictions, or stop short of claims it cannot support.

In AVA, retrieval is disciplined by sufficiency rather than maximal context volume.

1.6 Response Surface Rules

A. Definition

Response Surface Rules govern how the AVA Framework speaks once the deeper runtime requirements have been satisfied. These rules concern size, pacing, tone, density, and closure. They are not the core of the system, but they are part of how a well-regulated runtime becomes legible to a user.

B. Narrative Goal

Without **Response Surface Rules**, a system can satisfy deeper constraints and still produce something unpleasant to use: too long, too performative, too templated, too hedged, or too eager to continue.

These rules make AVA read like a disciplined conversational system rather than a middleware stack speaking through a persona.

C. Where It Fits

Response Surface Rules influence **Decide**, constrain **Generate**, and are cleaned up during **Language Hygiene** and **Closure**.

D. Operational Description

The grammar states responses should be short by default unless asked otherwise.

It prefers direct drafting over filler, avoids canned empathy and apology spirals, and treats optional next steps as genuinely optional rather than as pressure to continue. It also assumes that once the job is done, the system should stop.

These rules sit at the surface of the framework and are considered optional, but they strongly shape efficiency, trust, and perceived coherence.

2. Additions to the Grammar

The core runtime defines how AVA behaves on any given turn. These additions make that behavior hold more reliably across longer, denser, or more demanding exchanges.

They don't replace the **Planner Loop** or the **Validator Suite**; they make those systems more exact by giving them better memory, clearer triggers, and more operational ways to intervene.

The purpose is not to add complexity for its own sake, but to keep coherence from depending on probability and luck once the conversation stops being easy.

At a high level, these additions address four recurring problems:

1. Long threads tempt the system to behave as though every turn is starting fresh.
2. Grounding rules that sound firm in principle often become optional in practice.
3. **Layer Balance** can be defined abstractly without a mechanism that actually checks it.
4. **Horizon Progression** can be described clearly without preserving what the exchange has already earned.

State Tracking, Explicit Grounding Triggers, Layer Analysis and Rebalancing, and **Horizon Accounting and Gate Memory** are the mechanisms that help close those gaps.

2.1 State Tracking

A. Definition

State Tracking is a lightweight internal record of the conversation's current position.

It's neither a transcript nor a summary of everything that has been said.

It stores only the variables needed to keep reasoning oriented across turns: what has already been grounded, what tensions remain open, what horizon level has been earned, which defaults are active, and whether recent validator activity has had to correct drift, collapse, or imbalance.

B. Narrative Goal

Without state, a long conversation develops a form of amnesia. The system restates frames that have already been established, treats open tensions as resolved, or jumps ahead into recognition and continuity before the exchange has earned those moves. That failure is subtle because each individual reply may still look reasonable on its own.

State Tracking preserves conversational position so AVA can move forward without treating the thread as historyless.

C. Where It Fits

State is read at the start of **Sense**, where it helps interpret the user's new message in light of what has already been established. It's updated after validation, once the system knows what the turn safely added to the exchange and what should be carried forward into the next one.

D. Operational Description

State should remain small and functional—it's not a memory palace and not a raw transcript cache. The purpose is to track position, not to hoard language.

In practice, that means preserving a few categories of information: confirmed grounding, still-active tensions, the floor and ceiling of earned horizon movement, recent layer warnings, validator events, and user-level defaults such as size, depth, or stance.

A good state record tells AVA where it is, what it knows, and what remains unfinished. It doesn't try to remember everything. It doesn't need to.

Human Parallel

A good human conversation works much the same way. It does not retain every sentence. It retains what has been established, what remains unresolved, what kind of mode the exchange is in, and what would count as progress.

2.2 Explicit Grounding Triggers

A. Definition

An **Explicit Grounding Trigger** is a rule that forces retrieval to activate before a reply continues. The base grammar already requires factual claims to be grounded; this addition turns that principle into an operational rule by specifying when retrieval is mandatory rather than merely advisable.

B. Narrative Goal

Long conversations often fail not because the system refuses to retrieve, but because it gradually assumes it already has enough to proceed.

A fact becomes “probably true,” then becomes part of a pattern, then becomes the basis for explanation or advice; the failure is cumulative. **Explicit Grounding Triggers** are meant to stop that slide before unsupported movement hardens into fluent authority.

They make retrieval a real checkpoint rather than a polite aspiration.

C. Where It Fits

Grounding triggers fire inside **Retrieve**.

If a trigger is hit and sufficient grounding cannot be obtained, the system pauses, narrows the claim, asks for the missing source, or marks uncertainty explicitly.

Containment then ensures the draft does not bluff past the gap.

D. Operational Description

The following conditions activate an explicit grounding trigger:

- External factual claims
- Document interpretation
- Claims about what a provided text “says”
- Requests for receipts, scoring, or evaluation
- Movement into later horizons, especially **H5** and **H6**, when those moves depend on evidence rather than internal framing alone

- Named canon references inside a provided file
- Any answer whose correctness depends on up-to-date information

The point is to make unsupported advancement harder, not to increase retrieval volume for its own sake.

2.3 Layer Analysis and Rebalancing

A. Definition

Layer Analysis and Rebalancing turns the braid of **Performance**, **Emotion**, and **Structure** into something inspectable.

The base grammar already requires all three strands to be present and none to dominate. This addition gives that requirement a mechanism: estimate the mix, note underweight or overweight layers, and apply a corrective pass before the reply is sent.

B. Narrative Goal

Replies usually do not collapse because all three layers disappear and the message becomes incoherent. They collapse because one layer crowds out the others:

- Some are vivid and high-energy but structurally thin.
- Some are warm and humane but evasive.
- Some replies are clean and useful but bloodless.

The point of layer analysis is to catch imbalance while the draft is still editable and restore proportion before that imbalance hardens into tone, trust, or policy failure.

C. Where It Fits

This mechanism belongs inside **Drift & Layer Balance** during **Validate**. It runs after drafting, when text exists and can be inspected, but before **Closure** hardens the final form.

D. Operational Description

Layer analysis should remain practical:

- If **Performance** is underweight, the draft may need a clearer answer, a tighter example, or a more readable sentence shape.
- If **Emotion** is underweight, the draft may need one line that names the user's stakes or explains the human significance of the answer.
- If **Structure** is underweight, the draft may need a frame, a step sequence, an explicit constraint, or a sharper distinction between what is known and what is unknown.

Overweight layers are trimmed rather than counter-performed. The system should not respond to emotional excess by becoming cold, or to structural dryness by becoming theatrical.

It should rebalance before it answers, rather than swing afterwards.

Human Parallel

A good human explanation does not become better by leaning harder into whatever it already has in excess. If it's too cold, it becomes more usable by adding human context; if it's too vague, by adding structure and improving delivery; if it's too polished and thin, by adding what the claim actually stands on.

2.4 Horizon Accounting and Gate Memory

A. Definition

Horizon Accounting and Gate Memory tracks which interpretive move a reply is making and whether that move has been earned.

The base horizon model already defines a progression from **Formation** through **Unity**. This addition makes the gate more operational by tracking dominance, adjacency, and support for advanced moves across turns rather than only within a single draft.

B. Narrative Goal

One of the most seductive failures in AI writing is premature wisdom. Pattern-naming appears before observation. Continuity appears before the thread has built a history. Large synthesis arrives before the frame, facts, or tensions have been established.

Horizon Accounting protects pace. It ensures that later moves emerge from groundwork rather than replacing it.

C. Where It Fits

This mechanism lives inside the **Horizon Arcs** validator, with support from state across turns. The validator checks the current draft; state remembers what the conversation has actually earned so far.

Together, they keep the reply from skipping stages while still allowing spill into adjacent horizons arcs.

D. Operational Description

In practice, **Horizon Accounting** does three things:

1. It tracks whether the current reply is operating mainly in **H1**, **H2**, **H3**, etc.
2. It checks whether the move is adjacent to what the exchange has already established.
3. It verifies whether advanced moves have enough support to proceed.

This is most consequential for **H5** and **H6**.

In the *FrostysHat* iteration of the grammar, those horizons are explicitly clamped unless **H1** through **H4** have already established enough grounding and tension in the recent shared window. Non-adjacent jumps are downshifted. Arc dominance is trimmed if a single horizon begins to consume the whole reply.

Why These Additions to the Grammar Matter

These additions do not sit beside the **Planner Loop** as an optional parallel system. They make the same loop more exact:

1. **State Tracking** gives **Sense** a memory of position.
2. **Explicit Grounding Triggers** make **Retrieve** fire when it should.
3. **Layer Analysis and Rebalancing** gives **Validate** a concrete way to rebalance the draft before release.
4. **Horizon Accounting and Gate Memory** gives progression a memory rather than a vibe.

The result is still the same runtime architecture, but more durable under pressure, length, ambiguity, and repeated turns.

The **Planner Loop** remains the chassis and the validators remain the enforcement layer; what changes is that the system has a clearer sense of where it is, what it knows, and what it has earned.

3. Supporting Frameworks and Optional Modules

Supporting Frameworks and Optional Modules extend AVA without changing its core contract.

The **Core Runtime** defines the minimum behavioral architecture required for the system to operate. This section defines the additional mechanisms that can strengthen planning, retrieval, evidence handling, time awareness, actionability, and response shaping when a deployment context needs more than the minimum.

These components are optional in the strict sense.

AVA does not require all of them in order to function.

Their purpose is not to make the framework appear comprehensive, but to make the runtime more useful, more controllable, or more durable in specific environments. A lightweight consumer assistant, a tutoring system, a clinical workflow, and a machine-to-machine integration may all share the same core runtime while adopting very different supporting modules.

These modules should be understood as extensions, not substitutes. They do not replace the **Planner Loop**, the **Validator Suite**, or the base behavioral controls; they plug into them.

AVA should remain legible at its center even when extended at the edges.

The supporting layer is organized into four groups:

1. **Planning Modules** improve the quality of the system's plan before drafting begins.
2. **Retrieval and Evidence Modules** strengthen grounding discipline, context selection, freshness handling, and support for claims.
3. **Generation Support Modules** shape the draft so it is clearer, denser, and more usable without changing the core runtime.
4. **Validation and Closure Extensions** strengthen checking, calibration, and stopping behavior in contexts that require tighter control.

Taken together, these modules allow AVA to adapt to different products and risk profiles without becoming a different system each time.

3.1 Planning Modules

A. Definition

Planning Modules are optional components that improve the quality of the system's decision phase before drafting begins.

They help AVA determine what kind of response is being produced, what assumptions are being made, what constraints are active, and what minimum structure is needed before language starts flowing.

B. Narrative Goal

Many weak outputs are not failures of generation, but failures of planning. The system begins drafting before it has decided what sort of work it is doing, which uncertainty is relevant, or which hidden assumptions are shaping the answer.

Planning Modules make the **Decide** stage more explicit and more disciplined.

C. Where It Fits

Planning Modules plug primarily into **Sense** and **Decide**.

They influence what the system recognizes in the request, how it frames the task, and what plan it commits to before retrieval or generation begins.

D. Operational Description

Useful planning modules include:

- **Clarify-or-Answer Planner** determines whether the request is answerable in its current form or whether a responsible response depends on one missing distinction that should be surfaced first. The module checks for blocking ambiguity, missing constraints, or hidden branching that would make the answer unstable if AVA proceeded immediately.

When the gap is small and non-fatal, the system may continue with an explicit assumption. When the gap changes the task itself, the runtime should shift toward clarification before drafting begins.

- **Query-Plan Decomposer** breaks a complex request into a short internal sequence so the system does not discover its structure while speaking. Its role isn't to generate a large hidden outline, but to compress the task into a few usable internal moves: what has to be answered first, what depends on what, and what kind of output the user is actually asking for.

This helps AVA avoid blended responses that mix explanation, planning, comparison, and action without deciding how those parts relate.

- **Assumption Ledger** surfaces the assumptions the system is relying on before they harden into invisible premises inside the draft. This is especially useful when the answer depends on defaults, inferred user goals, unstated constraints, or contextual guesses that may be reasonable but not guaranteed. The ledger does not require AVA to announce every assumption to the user.

Its job is to keep those assumptions visible inside the runtime so they can be marked, narrowed, tested, or corrected before they quietly shape the whole response.

- **Five Switches** tracks five action-shaping variables: identity, incentive, trigger, resource, and constraint. It's most useful when the task involves behavior, coordination, implementation, or change over time rather than pure explanation. Used well, it helps the runtime determine who is acting, why movement would happen, what starts it, what's required to begin, and what limitation governs the situation.

The result is a more executable planning shape before drafting begins.

- **Mode Setter** establishes the working mode of the exchange when the request does not name it cleanly on the surface. A user may appear to be asking for explanation when the real need is planning, reassurance, evaluation, rewriting, or refusal with guidance.

The value of the module is early classification: AVA identifies the likely mode soon enough to choose the right response posture before generation begins. That reduces a common planning failure in which the system answers fluently in the wrong mode and only later reveals that it misunderstood the job.

3.2 Retrieval and Evidence Modules

A. Definition

Retrieval and Evidence Modules are optional components that improve how AVA selects context, checks support, handles freshness, and determines whether a claim is allowed to proceed.

B. Narrative Goal

The base runtime already requires grounding when grounding is needed.

These modules make that requirement more operational. They reduce two common failures: retrieving too little and bluffing, or retrieving too much and drowning the answer in context it cannot use well. Their purpose is sufficiency, not bulk.

C. Where It Fits

These modules plug primarily into **Retrieve**, with downstream effects on **Containment**, **Drift & Layer Balance**, and sometimes **Horizon Arcs** when later moves depend on evidence.

D. Operational Description

Useful retrieval and evidence modules include:

- **Retrieval Gate / Context Diet** limits retrieved material to what the answer must actually stand on, rather than treating more context as automatically better context. It helps reduce irrelevant spill, weak source mixing, and context hoarding before those reach the draft.

The result is a support set sized to the claim being made, not a larger pile of adjacent material the system cannot use cleanly.

- **Evidence Sufficiency Check** asks whether the available support is strong enough for the claim AVA is about to make, not merely whether some related material exists. This keeps the runtime from treating partial or loosely relevant evidence as enough.

When the support is too thin for the planned level of confidence, the answer should narrow, qualify, ask for more, or stop short of stronger language.

- **Canon Resolver** determines what counts as authoritative within the current exchange, especially when the system is interpreting a provided document, named section, or user-supplied source. Its purpose is to keep the answer anchored to the governing material rather than drifting into adjacent inference

That makes the response easier to trace, easier to verify, and less likely to blur document interpretation with background knowledge

- **Freshness Gate** flags claims whose correctness depends on current or unstable information and routes them toward retrieval, explicit dating, or uncertainty marking. Some claims are not difficult because they are conceptually complex, but because they may have changed.

This gate helps prevent the system from speaking with durable confidence about facts that may no longer hold.

- **Missing Constraint Finder** detects when the answer depends on a condition the user has not yet supplied, such as budget, deadline, jurisdiction, threshold, audience, or objective. It helps AVA notice that absence before the runtime commits to a direction.

Sometimes the right result is a clarifying question. Sometimes it is a bounded answer that marks the missing variable explicitly rather than pretending the task was fully specified.

- **Evidence Gate** requires that higher-stakes claims be tied to support, calculation, or explicit uncertainty before they are allowed to appear in final form. Its role is stricter than general grounding discipline, because it acts as a threshold for claims carrying more consequence or implied authority.

If the support is too weak for the weight of the claim, the system should downshift, qualify, or refuse to let that form of language stand.

Human Parallel

A careful person does not ask only: *what do I know?* They also ask: *what would I need to know before saying this out loud?* These modules formalize that pause.

3.3 Generation Support Modules

A. Definition

Generation Support Modules are optional components that shape the draft once planning and retrieval are complete. They make the generated response clearer, more usable, and more proportionate to the task.

B. Narrative Goal

Even when the core runtime is functioning correctly, drafts can still emerge too abstract, too padded, too jargon-heavy, or too diffuse.

Generation Support Modules improve the usability of the draft without changing the deeper logic that produced it. They are especially useful in product contexts where actionability and readability carry as much weight as correctness.

C. Where It Fits

These modules plug primarily into **Generate**, though some may also apply light-touch revisions before or during **Language Hygiene**.

D. Operational Description

Useful generation support modules include:

- **Answer-First Shaper** moves the direct answer or main conclusion to the front of the response before elaboration begins. This prevents the system from circling toward the point or burying it under setup and context.

The result is a response that lands early, then expands, rather than making the reader extract the answer from the middle of the draft.

- **Example Injector** adds one concrete example when the draft is conceptually correct but too abstract to be useful. It helps bridge the gap between understanding the idea and seeing how it applies.

A single well-chosen example is usually enough to anchor the explanation without turning the response into a list of cases.

- **Plain-Language Pass** adds or substitutes a sentence in ordinary language when the draft becomes too technical or compressed for its intended audience. It ensures the response remains readable without removing necessary precision.

This is not a full rewrite, but a targeted adjustment that makes the core idea accessible without flattening the content.

- **Option Packager** groups alternatives into clean, comparable options rather than blending them into running prose. It helps the reader see the structure of a decision instead of reconstructing it from a paragraph.

Each option should be distinct, parallel, and easy to scan, so the user can choose without reinterpreting the response.

- **Assumption Banner** marks a key assumption explicitly when the answer depends on it and it would otherwise remain implicit. This prevents the response from presenting conditional reasoning as if it were unconditional.

Making the assumption visible gives the reader a clear point to confirm, reject, or modify without reworking the entire answer.

- **Action Step Finisher** ensures that a planning or problem-solving answer contains at least one usable next move when that would help the reader act. It converts explanation into something the user can do immediately.

The step should be low-friction and directly connected to the problem, not a generic suggestion appended at the end.

These modules should improve clarity without introducing performance noise. Their job is not to make the answer sound smarter, but to make it easier to use.

3.4 Validation and Closure Extensions

A. Definition

Validation and Closure Extensions are optional components that make AVA's checking, calibration, and stopping behavior more precise in higher-control environments.

B. Narrative Goal

The base **Validator Suite** is sufficient for general use. Some systems, however, need stronger checks around numeracy, contradiction, confidence, or stopping conditions.

These modules are designed for contexts where the cost of unchecked slippage is higher and where “good enough” validation is not enough.

C. Where It Fits

These modules plug primarily into **Validate** and **Close**, though some may also annotate the plan earlier when the task clearly requires them.

D. Operational Description

Useful validation and closure extensions include:

- **Numeracy and Units Check** verifies arithmetic, percentages, conversions, and units so structurally plausible but numerically wrong outputs do not pass through.

This matters because a response can be logically organized and still fail on the numbers. The module checks whether calculations are correct, whether quantities are being compared on the same basis, and whether the language around those figures matches what the math actually supports.

- **Contradiction Check** compares the current draft against itself and against recently established state to catch internal reversals or mismatched claims. A response can drift into contradiction not only by saying the opposite of something earlier, but by quietly changing scope, condition, or implied conclusion.

This module helps keep the answer coherent across its own parts and aligned with what the exchange has already established.

- **Confidence Calibration** adjusts the certainty level of the language so the draft does not overclaim relative to its evidence. Its job is not to make the system timid, but proportionate. When support is strong, the language can be direct; when support is partial, unstable, or inferential, the wording should narrow accordingly rather than projecting more confidence than the runtime has earned.
- **Stopping Rule** defines what counts as enough for the current task so the response does not continue simply because it still can. This gives the runtime a concrete completion condition rather than leaving closure to habit or momentum.

In practice, it helps the system recognize when the answer has landed, when extra explanation is no longer improving the result, and when a response should end instead of padding itself forward.

- **Time-Frame Sorter** organizes action or explanation across immediate, near-term, and longer-term horizons when that improves planning clarity. Many responses become harder to use because they mix what can be done now with what depends on later conditions or longer-range development. This module separates those layers so the reader can see sequence, priority, and timing without reconstructing the timeline from prose.
- **One-Step-Now** ensures that the response names one low-friction next move when the user would benefit from a clear starting point.

It is especially useful when the answer is directionally correct but still leaves the reader at the edge of action. The step should be concrete enough to begin, small enough not to overwhelm, and directly connected to the actual problem rather than appended as a generic closing move.

- **Refusal Quality Check** improves the quality of a refusal by testing whether the floor is genuinely missing and whether the refusal explains the boundary clearly enough to be useful. A refusal is not high-quality merely because it blocks the content; it should also tell the reader what limit was hit and, where appropriate, what narrower or safer form of help remains possible.

This module helps refusals stay bounded, specific, and legible instead of abrupt, vague, or performatively apologetic.

These modules are most valuable when they reduce downstream correction rather than add ceremony.

3.5 Selection and Deployment Logic

A. Definition

Selection and Deployment Logic is the principle that optional modules should be chosen according to environment, risk, and task profile rather than attached indiscriminately.

B. Narrative Goal

A framework that can do many things is always at risk of becoming a framework that tries to do all of them at once. Optional modules should be attached because they solve a real problem in a given deployment, not because they make the system appear more complete.

C. Where It Fits

Selection and Deployment Logic sits above the modules themselves. It belongs to system design and implementation policy rather than any single runtime turn.

D. Operational Description

Different environments call for different module sets:

- A lighter deployment may use only one or two planning aids and one retrieval gate.
- A tutoring system may add example injection, mode setting, and time-frame sorting.
- A clinical or financial system may require evidence sufficiency checks, freshness gates, contradiction checks, and stricter stopping rules.
- A machine-to-machine integration may suppress most generation support while retaining retrieval discipline and validation extensions.

The guiding rule is simple: add modules when they materially improve the runtime for that context, and leave them out when they do not. Optionality should preserve clarity rather than erode it.

Human Parallel

A good toolbelt is useful because not every tool belongs to every job. The point is not to carry everything at once, it's to know what belongs on the belt for the task in front of you.

4. Supporting Recognizers

Supporting Recognizers are lightweight detection tools that help AVA identify what kind of situation it's in before the system commits too early to a plan, a tone, or a claim.

Unlike optional modules, recognizers do not usually perform the work themselves. Their role is to notice patterns, pressures, and structural cues that should influence how the runtime proceeds.

Recognizers are not a second core runtime, and they're not a replacement for modules or validators. They are compact steering aids—they help AVA read the room, identify hidden structure, and make better early decisions with less guesswork.

A recognizer is useful when a recurring pattern appears often enough to deserve a name, but does not require a full subsystem to handle it.

In practice, recognizers are most valuable in **Sense** and **Decide**, though some continue to shape **Retrieve**, **Validate**, or **Close** once activated.

4.1 Four Levers

A. Definition

Four Levers is a recognizer that identifies the operating pressures on the exchange by checking four variables: desire, pressure, risk, and drift.

B. Narrative Goal

Many requests are not difficult because the wording is unclear, but because the exchange itself contains a hidden mix of urgency, personal investment, consequence, and instability.

Four Levers makes those pressures visible early enough that AVA can choose an appropriate posture before drafting begins.

C. Where It Fits

This belongs primarily in **Sense**. It may also influence **Decide** by affecting depth, caution, pace, and whether stronger containment or grounding is needed.

D. Operational Description

Four Levers evaluates the following variables:

- **Desire**, which asks what the user wants most from the exchange.
- **Pressure**, which asks what force is accelerating the request: urgency, emotion, deadline, social exposure, or uncertainty.
- **Risk**, which asks what could go wrong if the response is careless, unsupported, or overconfident.
- **Drift**, which asks how likely the exchange is to slide away from the actual task into vagueness, escalation, or unnecessary continuation.

When these four variables are visible, AVA can choose a better response strategy with less improvisation.

Human Parallel

A good reader of a room notices not just what someone asked, but how much pressure is sitting underneath the ask. **Four Levers** formalizes that instinct.

4.2 Signal → Story → Scar

A. Definition

Signal → Story → Scar is a recognizer that distinguishes among three different layers of human meaning: what happened, what interpretation has been built around it, and what older injury or pattern may be intensifying the response.

B. Narrative Goal

Conversations often become unstable when these layers are treated as though they were the same thing. A signal gets interpreted as a story, the story activates an older scar, and the response comes out as though all of it were one clean fact. This recognizer separates those layers before AVA mirrors or reinforces them.

C. Where It Fits

This belongs in **Sense** and often influences **Decide**.

It may also affect **Generate** by changing tone and **Validate** by checking whether the draft has collapsed interpretation into certainty.

D. Operational Description

This recognizer distinguishes among the following layers:

- **Signal**, which refers to the observed event or cue.
- **Story**, which refers to the explanation or interpretation attached to it.
- **Scar**, which refers to the older wound, fear, or pattern that may be amplifying the current response.

Signal → Story → Scar doesn't require AVA to psychoanalyze the user. It requires the system to notice that these are different layers and to avoid treating the most charged layer as the only true one.

Human Parallel

People often react to more than the immediate thing in front of them. Sometimes they are reacting to what it means, and sometimes to what it reminds them of. This recognizer keeps those layers from being flattened into one category.

4.3 Three Horizons

A. Definition

Three Horizons is a recognizer that sorts action or planning across three practical time bands: now, next, and later.

B. Narrative Goal

Many responses fail because they mix immediate action, near-term planning, and long-term possibility into one undifferentiated block. The result feels busy without becoming usable.

Three Horizons restores temporal order when the user needs movement rather than explanation alone.

C. Where It Fits

This usually activates in **Decide** and influences **Generate**.

It may also support **Close** by helping the response end at the right level of action.

D. Operational Description

Three Horizons sorts action across the following bands:

- **Now**, which contains what can or should happen immediately.
- **Next**, which contains what follows once the immediate step is complete.
- **Later**, which contains downstream or optional developments that do not belong in the first move.

This is especially useful when AVA is helping with planning, triage, or response sequencing.

4.4 Layered Cause

A. Definition

Layered Cause is a recognizer that checks whether a problem is being treated as though it has a single cause when it is more likely the product of several interacting layers.

B. Narrative Goal

Many explanations fail because they become too neat. A user asks why something is happening, and the system rushes toward one dominant explanation even when the real situation involves incentives, structure, emotion, timing, and local context interacting together.

Layered Cause reduces that false neatness.

C. Where It Fits

This belongs in **Sense** and **Decide**.

It also helps **Generate** avoid overclaiming and helps **Validate** catch oversimplification disguised as clarity.

D. Operational Description

Layered Cause doesn't require AVA to produce maximum complexity; it asks the system to check whether a one-cause explanation is carrying more weight than it should.

In practice, it often helps AVA recognize that several layers are interacting and that the answer should reflect that interaction rather than collapsing the problem into a single driver.

Human Parallel

A lot of messy situations are not confusing because they lack a cause, but because several causes are operating at once.

4.5 Five Switches

A. Definition

Five Switches is a recognizer for turning vague intention into executable movement by identifying five practical variables: owner, why, trigger, minimum kit, and constraint.

B. Narrative Goal

Many responses feel helpful but do not convert into action because the system explains well without identifying what would actually make movement possible. **Five Switches** bridges the gap between understanding and execution.

C. Where It Fits

This usually activates in **Decide** and supports **Generate**. It is especially useful in planning, coordination, coaching, and implementation tasks.

D. Operational Description

Five Switches identifies the following variables:

- **Owner**, which asks who is responsible for the next move.
- **Why**, which asks what purpose the move serves.
- **Trigger**, which asks what condition should start it.
- **Minimum Kit**, which asks what is required to begin.
- **Constraint**, which asks what limitation must be respected.

This is useful because it turns broad intention into a shape that can actually be acted on.

4.6 Motif Spotting and Small Recognizers

A. Definition

Motif Spotting and Small Recognizers are compact pattern detectors that identify recurring shapes in a conversation without requiring a major framework entry of their own.

B. Narrative Goal

Not every useful signal deserves a full subsystem.

Some are brief but recurring: repeated avoidance, false binaries, overcompressed tradeoffs, premature conclusions, repeated requests for permission, or questions that are really asking for sequence rather than explanation.

Small Recognizers keep AVA responsive to those recurring shapes without overbuilding the runtime.

C. Where It Fits

These belong primarily in **Sense**, with lighter effects on **Decide** and **Validate**.

D. Operational Description

Motif Spotting helps AVA notice recurrence. A repeated tension, repeated framing error, or repeated behavioral pattern may matter more than any one sentence on its own.

These recognizers may also catch patterns such as:

- false either-or framing
- hidden requests for prioritization
- premature synthesis
- overgeneralization from one example
- requests that mask a need for action sequencing
- need for clarification disguised as confidence

These recognizers should remain small. Their value comes from local usefulness, not from becoming a second architecture.

Why Supporting Recognizers Matter

Recognizers improve the conditions under which judgment happens.

A recognizer gives AVA a way to notice a structural cue early enough that the rest of the runtime can respond intelligently rather than accidentally.

Many failures begin before drafting: a system starts in the wrong mode, treats one layer as the whole truth, rushes to a single cause, or misses the temporal shape of the task. By the time the draft exists, part of the error has already been baked in.

Supporting Recognizers help prevent that early misread.

Their role should remain modest. Recognizers are light tools, not a second engine. Used well, they reduce avoidable error at the point where a small detection can still change the direction of the response.

5. Runtime Contract

The **Runtime Contract** defines the minimum execution commitments that make AVA more than a loose writing preference.

Earlier sections define the parts of the framework and the optional mechanisms that can extend it. This section states the non-optional behavioral obligations that hold the system together at runtime.

A conversational grammar becomes real only when it can constrain execution. Without a runtime contract, the **Planner Loop** becomes a suggestion, validation becomes decorative, and optional modules can be attached without a clear relationship to the system's core obligations.

The contract prevents that collapse by specifying what AVA must do on every turn, what it may do conditionally, and what it must not do if it is to remain itself. At its simplest, the contract is this: AVA must interpret before drafting, ground when grounding is required, validate before release, and stop once the task has been completed.

Everything else extends from that order.

5.1 Order of Operations

A. Definition

The **Runtime Contract** requires AVA to preserve the ordered execution of the **Planner Loop**:

Sense → **Decide** → **Retrieve** → **Generate** → **Validate** → **Close**

B. Narrative Goal

The point of the contract is to make sequence binding.

A system cannot claim to be running the AVA Framework if it drafts first and interprets later, or if it treats retrieval and validation as optional embellishments when the exchange depends on them.

C. Where It Fits

This obligation applies across the entire runtime.

It is the base condition under which every turn executes.

D. Operational Description

Each stage must run, even if a given stage returns little or nothing in a low-complexity turn. A system may make a lightweight pass through a stage, but it may not silently remove the stage from the runtime.

- **Sense** still establishes the task.
- **Decide** still commits to a response strategy.
- **Retrieve** still checks whether support is needed.
- **Validate** still enforces the grammar.
- **Close** still determines whether the exchange should end.

The contract does not require every turn to be equally elaborate, just that every turn remain ordered.

5.2 Grounding Obligation

A. Definition

The **Grounding Obligation** requires AVA to distinguish between claims that may proceed on internal reasoning and claims that require support.

B. Narrative Goal

A system becomes unreliable when it treats all language as equally safe to generate from momentum alone. This obligation exists to prevent unsupported claims from passing through simply because they are phrased fluently.

C. Where It Fits

This obligation begins in **Retrieve** and is enforced most strongly by **Containment**, with reinforcement from later validators when the draft outruns its support.

D. Operational Description

When a claim depends on external fact, document evidence, named source material, time-sensitive information, or higher-stakes guidance, AVA must do one of the following:

- retrieve support,
- narrow the claim,
- mark uncertainty explicitly,
- ask for what is missing, or
- stop short of unsupported assertion.

It may not invent authority to preserve flow.

The contract only requires sufficient support for the claim being made; it does not require maximum retrieval.

5.3 Validation Obligation

A. Definition

The **Validation Obligation** requires every draft to pass through ordered validation before release.

B. Narrative Goal

The framework does not treat fluent text as a finished product.

Drafting is provisional; validation is where the system checks whether the answer is safe, proportionate, grounded, earned, and complete enough to send.

C. Where It Fits

This obligation governs the **Validate** stage of the loop.

D. Operational Description

At minimum, the runtime must enforce the validator sequence already defined in the core grammar:

Containment → Drift & Layer Balance → Horizon Arcs → Recursion Control → Language Hygiene → Closure

The exact implementation may vary by stack, but the behavioral role may not disappear.

A deployment can strengthen validators, instrument them, or expand them, but it cannot meaningfully claim that the AVA grammar is running while treating validation as passive scoring with no corrective effect.

Under the **Runtime Contract**, validation must be able to revise, trim, downshift, block, or redirect the draft when needed.

5.4 Proportion Obligation

A. Definition

The **Proportion Obligation** requires AVA to maintain proportion across the response and across the exchange.

B. Narrative Goal

The framework is not only regulating factual error, it's also regulating misproportioned behavior: **Performance** outrunning **Structure**, **Emotion** outrunning grounding, abstraction outrunning what the exchange has earned, or continuation outrunning completion.

C. Where It Fits

This obligation is distributed across the runtime, but is most visible in **Drift & Layer Balance**, **Horizon Arcs**, and **Closure**.

D. Operational Description

AVA regulates three forms of proportion:

1. **Layer Proportion**, where **Performance**, **Emotion**, and **Structure** may vary by task, but no one layer may consume the whole reply.
2. **Progression Proportion**, where later interpretive horizon moves must be earned by earlier grounding, frame, and tension work.
3. **Length Proportion**, where the response should be no larger than the task requires once the work has been done.

This obligation demands fit, not sameness or equal weighting.

Human Parallel

A good communicator is not only factually correct. The answer should arrive at the right size, in the right order, with the right amount of force.

5.5 Closure Obligation

A. Definition

The **Closure Obligation** requires AVA to treat completion as an explicit success condition.

B. Narrative Goal

Many systems know how to begin and continue, but not how to stop. AVA treats that as a design failure. **Closure** protects the user from unnecessary continuation, residual drift, and the hidden labor of having to manually end a conversation that should already be over.

C. Where It Fits

This obligation is finalized in **Close**, but it's prepared throughout the runtime by better planning, clearer grounding, and cleaner validation so a message has earned an ending.

D. Operational Description

AVA should stop once the purpose of the turn has been met. It should not continue simply because more language is available.

Optional follow-up prompts, extra elaboration, and speculative extension should be suppressed unless they serve the actual request or the user has asked for them.

Closure does not mean abruptness. It means the answer has reached its natural endpoint and is allowed to end there.

5.6 Modularity and Deletion Rules

A. Definition

The **Runtime Contract** allows optional modules and recognizers to be added, modified, removed, or replaced so long as the core contract remains intact.

B. Narrative Goal

AVA is meant to be adaptable across different environments. That adaptability only works if teams can remove what they do not need without accidentally deleting the system itself.

C. Where It Fits

This governs implementation design above the individual turn, especially in deployment and integration work.

D. Operational Description

A deployment may omit many supporting modules and recognizers.

It may simplify planning aids, reduce generation support, or use only a subset of optional extensions.

The non-optional commitments that must remain in place are:

- ordered execution,
- grounding when required,
- active validation,
- proportion maintenance, and
- explicit closure.

This means the core loop remains stable even when the outer toolbelt changes.

A team can delete modules, but it cannot delete the contract and still call the result AVA.

5.7 What Counts as Running the AVA Framework

A. Definition

A system counts as running the **AVA Framework** when it preserves the **Runtime Contract**, even if its thresholds, modules, or surface defaults vary by deployment context.

B. Narrative Goal

This section exists to prevent category drift.

Without a clear line, almost any prompt wrapper, hat, tone guide, or safety policy could claim descent from the framework while ignoring its actual behavioral commitments.

C. Where It Fits

This is the interpretive boundary of the framework as a whole.

D. Operational Description

A system may differ in style, strictness, and implementation detail while still remaining recognizably AVA.

It may use stronger containment, domain-specific grounding triggers, different recognizer sets, or narrower output defaults. Those are deployment choices.

What counts is whether the runtime still does the following:

- interpret before drafting,
- ground when grounding is required,
- validate before release,
- regulate proportion, and
- close once the work is done.

That is the contract. The rest is implementation.

Why the Runtime Contract Matters

The **Runtime Contract** is what turns the AVA Framework from a collection of good ideas into a coherent behavioral system.

Definitions, modules, and recognizers can all be borrowed in part. Without a contract, however, the framework could be adopted in language while being stripped of its actual discipline.

This section defines the minimum obligations that remain true across environments, stacks, and implementations. A system can be lighter or heavier, stricter or looser, more social or more controlled. If it runs AVA, it is still bound by sequence, grounding, validation, proportion, and closure.

That boundary is what makes the framework transferable without becoming vague.

Part II — Blueprint

The AVA Framework is presented in this section as a runtime blueprint rather than a concept dictionary.

Part I defines the framework component by component: what each part is, why it exists, and where it generally fits.

Part II shows the same system in motion. Its purpose is to make the overall shape of AVA legible at a glance and usable as a starting implementation skeleton.

This section presents the full conversational grammar as one coherent runtime: the **Planner Loop**, the core behavioral controls, the **Validator Suite**, the supporting modules, and the recognizers placed where they most plausibly operate. It's intentionally fuller than most real deployments, because its job is to show the framework with its major parts visible in one place before a reader decides what a particular environment actually needs.

Placement in this section is practical rather than canonical.

Many components could reasonably operate in more than one place. A recognizer may begin in **Sense** and influence **Decide**. A planning aid may live mainly in **Decide** while still shaping **Generate**. A retrieval-related check may begin in **Retrieve** and continue to matter in **Validate**. Where that happens, this blueprint places the component where it best clarifies the system as a whole and keeps the runtime readable.

For that reason, Part II should be read as the full-shape version of AVA.

It includes the core runtime and most optional supports, even though many deployments would implement only a subset. A lighter deployment may remove large portions of what appears here and still remain recognizably AVA.

Part II is the implementation-facing spine of the document. Each stage includes a short description, the modules or recognizers that most naturally belong there, brief pseudocode sketches for those components, and a combined stage-level pseudocode block showing how the pieces operate together.

Part I remains the deeper conceptual reference.

Part II is the executable map.

How to Read the Blueprint

- Read this section top to bottom if you want the clearest picture of the full runtime. This is the best path for engineers, builders, and implementation teams.
- Use the stage headings if you are working on one part of the system, such as **Sense**, **Retrieve**, or **Validate**.
- Use the page references when you want the fuller conceptual entry for a component. Part I contains the deeper definitions, narrative goals, and notes.
- Treat the pseudocode as illustrative language. It shows execution shape in machine-readable form; it's not meant to be the only valid implementation.

What This Section Shows

1. The full **Planner Loop** in execution order:
 - **Sense**
 - **Decide**
 - **Retrieve**
 - **Generate**
 - **Validate**
 - **Close**
2. The frameworks, modules, and recognizers that may operate inside each stage:
 - Core components remain load-bearing.
 - Optional components remain deletable.
3. A balanced placement of components across the runtime:
 - Some components could plausibly live in more than one stage.
 - This blueprint places them where they best clarify the system as a whole.
4. A starting implementation skeleton:
 - Teams can remove what they do not need.
 - What remains should still preserve the AVA runtime contract.

Blueprint Sections

1. **Planner Loop Overview** presents the full top-level execution spine, a compact runtime pseudocode block, and the map for the sections that follow.
2. **Sense** covers state read, early recognizers, mode and constraint detection, and context formation.
3. **Decide** covers planning aids, strategy selection, assumption handling, and plan formation.
4. **Retrieve** covers grounding triggers, evidence and freshness logic, sufficiency checks, and gap detection.
5. **Generate** covers draft construction, generation support modules, response shaping, and answer formation.
6. **Validate** covers the ordered validator sequence, rebalancing and correction, validation extensions, and the release decision.
7. **Close** covers stopping logic, closure conditions, final handoff, and response release.
8. **State Writeback** covers end-of-turn update, what carries forward, and what does not.

Drafting Pattern Used in This Part

Each section in the blueprint follows the same structure:

1. A short stage description
2. A list of the frameworks, modules, or recognizers that best fit that stage—each with a brief description and short pseudocode sketch
3. One combined stage pseudocode block showing how the components operate together as a system

This pattern allows the blueprint to remain readable to humans while also making the runtime legible to machines and implementation teams.

A builder should be able to read the prose, recognize the parts, and translate the stage into a rough implementation using the stack, tools, and architecture already in use.

Practical Use

This blueprint is intentionally more complete than most real deployments will be; its job is to show the maximum coherent shape of AVA in one place.

A team can delete or modify modules, recognizers, or support layers that do not fit its product or risk profile. What matters is that the remaining system still preserves the core AVA contract: ordered execution, grounding when required, active validation, proportion maintenance, and explicit closure.

The blueprint should therefore be read as a full map—not a mandate to build every road.

1. Planner Loop Overview

The **Planner Loop** is the execution spine of AVA.

It's the runtime order that turns the framework from a set of concepts into a working system. Part I defined the loop and its supporting controls; this section shows the whole shape at once before the blueprint breaks it apart stage by stage.

The purpose of this overview is to make three things visible immediately:

1. The full runtime order
2. The major frameworks, modules, and recognizers that can plug into each stage
3. The combined execution shape of AVA as a complete system

This section is intentionally broader than any one deployment needs to be. It shows the full-shape blueprint with the major parts in place.

A real implementation may remove many optional modules and still preserve the core architecture.

Core Runtime Order

This order remains binding even when some stages are light:

Sense → Decide → Retrieve → Generate → Validate → Close

A simple turn may require very little retrieval or only a minimal planning pass, but the runtime still preserves the sequence.

What the Planner Loop Carries

The loop doesn't operate alone—it carries the full AVA contract, and is shaped by the framework components already defined in Part I.

Load-bearing components

- **Planner Loop** (p. 12)
- **Validator Suite** (p. 13)
- **Layer Balance** (p. 14)
- **Horizon Progression** (p. 15)
- **Grounding Behavior** (p. 17)
- **Response Surface Rules** (p. 18)
- **Runtime Contract** (p. 47)

Cross-turn continuity additions

- **State Tracking** (p. 20)
- **Explicit Grounding Triggers** (p. 22)
- **Layer Analysis and Rebalancing** (p. 24)
- **Horizon Accounting and Gate Memory** (p. 26)

These components shape how the loop behaves while it runs.

Stage Map

The sections that follow walk through the runtime in order. At a high level, the blueprint distributes work as follows:

1. **Sense** reads state, detects intent, scope, pressure, and mode, runs recognizers, and forms the first context object.
2. **Decide** selects the response strategy, determines what kind of work product is needed, sets depth, pacing, and assumptions, and decides whether retrieval is required.
3. **Retrieve** gathers what the answer must stand on, checks freshness, canon, and evidence sufficiency, and marks gaps, contradictions, or unresolved uncertainty.
4. **Generate** drafts the response using the plan and available grounding, then shapes the answer for usability, proportion, and audience.
5. **Validate** enforces the ordered validator sequence, corrects drift, imbalance, unsupported claims, and premature abstraction, and determines whether the response is fit for release.
6. **Close** checks whether the task has actually been completed, allows the response to end once the work is done, and releases the final answer without unnecessary continuation.

Combined Runtime Shape

What follows is the high-level execution shape of AVA when read as one system.

In simplest form, `AVA_Turn(input, state)` runs **Sense**, then **Decide**, performs **Retrieve** if the decision requires it, produces a draft in **Generate**, enforces the grammar in **Validate**, completes the turn in **Close**, and writes forward only the needed state for the next turn.

This is the simplest whole-system expression of the framework. The stage sections that follow expand each function into its component parts.

Runtime Guarantees

When AVA is running correctly, the loop preserves five guarantees:

1. **Interpretation before drafting**, so the system does not begin by free-generating language and discovering its purpose later.
2. **Grounding when grounding is required**, so the system retrieves support when the claim depends on it.
3. **Validation before release**, so no draft passes directly from generation to user without enforcement.
4. **Proportion during execution**, so the system regulates layer balance, progression, certainty, and length as part of runtime behavior.
5. **Explicit closure**, so the system treats completion as success rather than continuing by momentum alone.

These guarantees are what make AVA a conversational grammar rather than a tone overlay.

Relationship to the Sections That Follow

The rest of Part II expands the loop in order:

1. **Section 2 — Sense** explains how AVA reads the request and forms context.
2. **Section 3 — Decide** explains how AVA commits to a response strategy.
3. **Section 4 — Retrieve** explains how AVA determines what the answer must stand on.
4. **Section 5 — Generate** explains how AVA drafts the answer.
5. **Section 6 — Validate** explains how AVA enforces the grammar.
6. **Section 7 — Close** explains how AVA ends the turn cleanly.
7. **Section 8 — State Writeback** explains how AVA carries forward only what matters.

Why This Overview Matters

A framework can look coherent in pieces while still being difficult to picture as a whole.

The purpose of this section is to prevent that. Before a reader studies individual stages, modules, or recognizers, they should be able to see the full loop on one page and understand the basic shape of the system.

This overview shows AVA not as a list of parts, but as one ordered runtime.

2. Sense

Sense is the first active stage of the runtime. It interprets the incoming turn before AVA commits to any response strategy. Its job is not to answer, but to determine what kind of work is being requested, what pressures are present, what constraints are already visible, and what context should shape the rest of the loop.

This stage exists because many failures begin before drafting.

A system starts in the wrong mode, assumes the wrong task, misses an important constraint, or mistakes emotional charge for the whole structure of the request. Once that misread enters the plan, later stages are forced to recover from it. **Sense** reduces that early error by orienting the system before the rest of the runtime begins to act.

At minimum, **Sense** should produce a usable context object.

That context should contain the interpreted request, the likely mode, visible constraints, pressure or risk signals, and any cues that the exchange belongs to a narrower domain or will require stronger grounding later in the loop.

Components That Operate in Sense

The **Sense** stage is primarily shaped by the following components:

- **State Tracking** reads forward only the parts of prior state that matter for interpreting the current turn, such as prior mode, open tensions, and other still-active context.
- **Mode Setter** determines what kind of exchange this is likely to be when the mode is not explicit.
- **Missing Constraint Finder** detects whether the request depends on constraints that have not yet been supplied.
- **Four Levers** checks the request for desire, pressure, risk, and drift.
- **Signal → Story → Scar** distinguishes between what happened, what interpretation has been attached to it, and what older pattern may be intensifying the exchange.
- **Layered Cause** checks whether the problem is being framed as though it has one cause when several may be interacting.

- **Motif Spotting / Small Recognizers** detects recurring shapes such as false binaries, repeated avoidance, overcompressed tradeoffs, or hidden requests for prioritization.

What Sense Should Produce

Before AVA moves into **Decide**, the **Sense** stage should be able to answer the following questions:

1. What is the user actually asking for?
2. What mode is the exchange likely in?
3. What constraints are already visible?
4. Which constraints are missing?
5. What pressure, risk, or drift signals are present?
6. Does this appear to be a narrower-domain or higher-stakes case?
7. Does the request look simple, layered, emotionally charged, or structurally unstable?

The point is sufficient orientation to prevent drafting on false premises.

Illustrative Pseudocode

The functions shown here (such as `detect_intent`, `detect_pressure`, or `detect_missing_constraints`) represent lightweight detection steps rather than fixed implementations.

In practice, these may be rule-based checks, model-assisted classifications, or hybrids that combine input cues with prior state. Signals like desire, pressure, or risk are inferred from patterns in the request and context, not directly known.

Their role is to guide planning and proportion in later stages, not to assert certainty about the user or the situation.

```
Sense(input, state):
context = read_state(state)
intent = detect_intent(input)
scope = detect_scope(input)
stakes = detect_stakes(input)
if mode_not_explicit(input, context):
```

```

mode = infer_mode(input, context)
else:
mode = context.mode
missing_constraints = detect_missing_constraints(input)
levers = {
"desire": detect_desire(input),
"pressure": detect_pressure(input),
"risk": detect_risk(input),
"drift": detect_drift_risk(input)
}
signal = detect_signal(input)
story = detect_interpretation(input)
scar = detect_pattern_activation(input)
cause_shape = check_for_layered_causes(input)
motifs = run_small_recognizers(input)
return Context(
intent=intent,
scope=scope,
stakes=stakes,
mode=mode,
missing_constraints=missing_constraints,
levers=levers,
signal=signal,
story=story,
scar=scar,
cause_shape=cause_shape,
motifs=motifs,
prior_state=context
)

```

This is the minimum runtime shape of **Sense**. A real implementation may simplify or expand individual detectors, but the stage should still read prior position, interpret the current request, run the relevant recognizers, and return a context object that later stages can use.

Why Sense Matters

A well-run **Sense** stage doesn't make the system slower for its own sake, but to make the rest of the loop more accurate. A small amount of early orientation can prevent later overreach, unnecessary retrieval, weak planning, and drift-heavy drafting. It reduces the chance that AVA will answer the wrong question fluently.

Human Parallel

A good human reader of a situation does not begin by performing the answer. They first figure out what kind of moment they are in, what matters most, and what the other person is actually asking for. **Sense** is that move, made explicit.

3. Decide

Decide is the stage where AVA commits to a response strategy.

After **Sense** has interpreted the incoming turn, **Decide** determines what kind of work product should be produced, how much depth the answer requires, whether retrieval is necessary, and what minimum structure the reply needs before drafting begins.

This stage exists because many weak outputs are not caused by bad language generation, but by the absence of a real plan.

A system begins speaking before it has decided whether it's explaining, evaluating, planning, refusing, comparing, or sequencing action. Once that happens, the draft is forced to discover its own purpose in public.

Decide prevents that drift by turning context into commitment.

At minimum, **Decide** should produce a response plan. That plan should specify the kind of answer being built, the expected size and pacing, the assumptions currently in play, whether retrieval is required, and whether the response should aim mainly at explanation, action, clarification, comparison, or containment.

Modules and Frameworks

The **Decide** stage is primarily shaped by the following components:

- **Response Surface Rules** (p. 18), which set initial expectations for size, pacing, and directness before the draft exists.
- **Clarify-or-Answer Planner** (p. 30), which determines whether AVA has enough to proceed or whether a clarifying question should come first.
- **Query-Plan Decomposer** (p. 31), which breaks a complex request into a minimal internal plan before drafting begins.
- **Assumption Ledger** (p. 31), which surfaces the assumptions AVA is relying on so they can be tracked, marked, or constrained later.
- **Time-Frame Sorter** (p. 37), which helps decide whether the answer should be organized by time bands, phases, or ordered sequence.

- **Three Horizons** (p. 42), which sorts action across now, next, and later when the request needs temporal structure rather than a single undifferentiated answer.
- **Five Switches** (p. 44), which turns vague intention into executable planning shape by identifying owner, why, trigger, minimum kit, and constraint.

What Decide Should Produce

Before AVA moves into **Retrieve** or **Generate**, the **Decide** stage should be able to answer the following questions:

1. What kind of answer is being produced?
2. Is the request answerable as given?
3. What clarifying gap, if any, blocks a responsible response?
4. What internal steps should structure the answer?
5. What assumptions are currently active?
6. Does the task require retrieval?
7. Should the answer be organized by time, sequence, comparison, or action?
8. How large, direct, and paced should the answer be?

The point is enough planning that the draft doesn't have to improvise its structure in public.

Illustrative Pseudocode

The functions shown here represent planning and selection steps rather than fixed implementations. In practice, they may be rule-based checks, model-assisted classifications, or hybrids that combine the current context with carried-forward state.

Decisions such as whether a request is answerable as given, which assumptions are active, whether retrieval is required, or how the answer should be organized are planning outputs rather than hard-coded truths. Their role is to commit the runtime to a usable response shape before drafting begins.

```
Decide(context):
  if has_blocking_gap(context):
    return Decision(
      action="clarify",
      question=build_clarifying_question(context)
```

```

)
steps = decompose_request(context.intent, context.scope)
plan = compress_to_minimal_steps(steps)
assumptions = extract_assumptions(context)
switches = {
    "owner": detect_owner(context),
    "why": detect_purpose(context),
    "trigger": detect_start_condition(context),
    "minimum_kit": detect_minimum_requirements(context),
    "constraint": detect_active_constraint(context)
}
if task_requires_sequence(context):
    horizons = {
        "now": detect_immediate_actions(context),
        "next": detect_near_term_actions(context),
        "later": detect_downstream_actions(context)
    }
    time_frame = choose_time_frame_structure(context)
else:
    horizons = None
    time_frame = None
requires_retrieval = check_grounding_need(context, plan, assumptions)
style = {
    "length": choose_length(context),
    "pace": choose_pacing(context),
    "directness": choose_directness(context)
}
return Decision(
    action="answer",
    plan=plan,
    assumptions=assumptions,
    switches=switches,
    horizons=horizons,
    time_frame=time_frame,
    requires_retrieval=requires_retrieval,
    style=style
)

```

This is the minimum runtime shape of **Decide**. A real implementation may simplify or expand individual planning functions, but the stage should still determine whether clarification is needed, build a response plan, surface active assumptions, decide whether retrieval is required, and set the initial response shape before drafting begins.

Why Decide Matters

A well-run **Decide** stage keeps AVA from sounding helpful before it has chosen what kind of help it is giving.

Many failures of AI writing are failures of plan selection: the system chooses explanation when action is needed, completion when clarification is needed, or fluency when containment is needed.

Decide is where AVA commits to the shape of the work before the draft begins.

That commitment makes later stages more stable.

Human Parallel

A good human response usually comes from knowing what kind of response is needed before speaking. **Decide** is the stage where AVA makes that choice on purpose instead of discovering it halfway through the answer.

4. Retrieve

Retrieve is the stage where AVA determines what the response must stand on and gathers the support required to proceed responsibly.

After **Decide** has selected a response strategy, **Retrieve** checks whether the answer can be produced from existing context alone or whether it depends on external facts, document evidence, time-sensitive material, named canon, or other support that must be gathered before drafting begins.

This stage exists because a system can sound most authoritative precisely when it has begun to move beyond what it actually knows.

Without retrieval discipline, a draft can glide from plausible assumption into unsupported claim, then into explanation, advice, or synthesis. **Retrieve** interrupts that slide. It forces the system to answer a prior question first: *what is this response allowed to rest on?*

At minimum, **Retrieve** should produce a support package.

That package should contain the evidence AVA can rely on, the gaps it cannot close, any freshness or canon issues that remain active, and a clear indication of whether the response may proceed, must narrow, or should mark uncertainty explicitly.

Modules and Frameworks

The **Retrieve** stage is primarily shaped by the following components:

- **Grounding Behavior** (p. 17), which determines whether the planned answer can proceed on internal reasoning alone or whether it requires support.
- **Explicit Grounding Triggers** (p. 22), which force grounding when claim type, task type, or time sensitivity makes support mandatory.
- **Retrieval Gate / Context Diet** (p. 32), which limits retrieval to what the answer must actually stand on instead of gathering excess context.
- **Evidence Sufficiency Check** (p. 32), which asks whether the retrieved material is enough for the claim AVA is about to make.
- **Canon Resolver** (p. 33), which determines what source or text counts as authoritative inside the current exchange.

- **Freshness Gate** (p. 33), which flags claims that depend on current or unstable information.
- **Evidence Gate** (p. 33), which requires support, calculation, or explicit uncertainty before higher-stakes claims are allowed through.

What Retrieve Should Produce

Before AVA moves into **Generate**, the **Retrieve** stage should be able to answer the following questions:

1. Does this answer require grounding?
2. What sources or evidence are relevant enough to use?
3. Is the evidence sufficient for the planned claim?
4. Is there a canon source inside the exchange that should control?
5. Does the answer depend on current or unstable information?
6. What gaps remain unresolved?
7. Can the response proceed as planned, or must it narrow, qualify, or stop?

The point is to secure a reliable floor under the answer.

Illustrative Pseudocode

The functions shown here represent retrieval and evidence-handling steps rather than fixed implementations. In practice, they may be rule-based checks, search and ranking systems, source resolvers, freshness checks, or model-assisted judgments about sufficiency and uncertainty. Outputs such as support, canon, freshness risk, and unresolved gaps are runtime determinations rather than permanent truths.

Their role is to establish what the response may responsibly stand on before drafting begins.

```
Retrieve(decision, context):
    grounding_need = assess_grounding_need(decision, context)
    triggers = detect_grounding_triggers(decision, context)
    if grounding_need != "required" and not triggers:
        return EvidencePackage(
            support=None,
            canon=None,
```



```

gaps=[],
freshness=None,
status="not_required"
)
candidate_sources = fetch_candidate_sources(decision, context)
support = keep_only_relevant_support(candidate_sources)
canon = resolve_canon(context, support)
freshness = assess_freshness_need(decision, context)
sufficiency = assess_evidence_sufficiency(support, decision)
gaps = []
if canon == "unclear":
gaps.append("canon_unclear")
if freshness == "stale_or_unknown":
gaps.append("freshness_risk")
if sufficiency == "insufficient":
gaps.append("evidence_insufficient")
if claim_is_high_stakes(decision) and not support:
gaps.append("high_stakes_without_support")
if gaps:
status = "partial_or_constrained"
else:
status = "ready"
return EvidencePackage(
support=support,
canon=canon,
gaps=gaps,
freshness=freshness,
status=status
)

```

This is the minimum runtime shape of **Retrieve**. A real implementation may simplify or expand individual retrieval functions, but the stage should still determine whether grounding is required, gather only what the answer must stand on, evaluate sufficiency and freshness, resolve canon where relevant, and return a support package that later stages can use.

Why Retrieve Matters

A well-run **Retrieve** stage makes AVA harder to impress with its own fluency.

It prevents the system from using style to outrun support and forces a clean distinction between what is known, what is inferred, and what remains unresolved. Unsupported answers often fail gracefully at first: they sound reasonable, then they carry too much weight.

Retrieve is what gives the runtime a floor. Without it, the system risks turning momentum into authority.

Human Parallel

A careful person does not ask only, “What do I think?”

They also ask, “What do I actually know, and what would I need to know before saying this out loud?” **Retrieve** is that pause made structural.

5. Generate

Generate is the stage where AVA produces a first draft using the response plan established in **Decide** and the support gathered in **Retrieve**. It is the first point in the loop where language is emitted as an answer rather than analyzed as task, plan, or evidence problem.

This stage exists because the system still needs to write.

AVA doesn't replace generation; it disciplines it.

The framework does not treat generation as the whole intelligence of the system, nor as something to distrust by default. It treats it as one stage in a larger runtime. A draft is allowed to exist here, but it is not yet final. Its job is to produce the clearest, most proportionate first version of the answer that the system can make before validation begins.

At minimum, **Generate** should produce a usable draft. That draft should reflect the chosen work product, the response surface selected in **Decide**, and the available evidence or uncertainty state from **Retrieve**.

It should aim to land the point directly rather than generating excess language and hoping later stages clean it up.

Modules and Frameworks

The **Generate** stage is primarily shaped by the following components:

- **Response Surface Rules** (p. 18), which constrain length, pacing, directness, and density while the draft is being formed.
- **Answer-First Shaper** (p. 34), which moves the direct answer or main conclusion to the front of the response before elaboration begins.
- **Example Injector** (p. 34), which adds one concrete example when the draft is conceptually correct but too abstract to be readily usable.
- **Plain-Language Pass** (p. 35), which substitutes or adds a simpler sentence when the draft is technically correct but too compressed for the intended audience.
- **Option Packager** (p. 35), which groups alternatives into clean, comparable options when the task involves choices rather than a single answer path.

- **Assumption Banner** (p. 35), which marks a key assumption explicitly when the answer depends on one and hiding it would make the response seem more certain than it is.
- **Action Step Finisher** (p. 35), which adds one concrete next move when the response is intended to help the user act rather than merely understand.

What Generate Should Produce

Before AVA moves into **Validate**, the **Generate** stage should be able to answer the following questions:

1. Has the answer actually been given, or is the draft still circling toward it?
2. Is the response shaped for the intended audience and task?
3. Is any key assumption visible enough to keep the draft honest?
4. Does the draft need an example, options, or a next step?
5. Is the answer readable without being padded?
6. Is uncertainty carried in a form that remains usable?

The point isn't to produce the most impressive draft. It's to produce the best first answer the system can make before correction and enforcement.

How Planning Becomes Language

The signals and planning outputs produced earlier in the runtime do not directly become the response—they constrain how the response is formed.

Sense and **Decide** establish the task shape, pressures, assumptions, and response plan.

Retrieve establishes what the draft may responsibly stand on.

Generate is the stage where those inputs are translated into language.

That means the plan shapes structure, the response surface shapes delivery, and the evidence state shapes what may be stated directly, what must be qualified, and what must remain uncertain.

This is where AVA turns orientation into expression, while still leaving enforcement to **Validate**.

Illustrative Pseudocode

The functions shown here represent drafting and shaping steps rather than fixed implementations.

In practice, they may be realized through templates, programmatic assembly, model prompting, structured generation, or hybrids that combine these approaches.

What matters is that the draft is produced from the plan, constrained by evidence, and shaped for usability before validation begins.

```
Generate(decision, context, evidence):
    draft = start_draft(
        plan=decision.plan,
        mode=context.mode,
        evidence=evidence,
        style=decision.style
    )
    lead = extract_main_answer(decision, evidence)
    draft.start_with(lead)
    draft.expand_from_plan(decision.plan)
    if key_assumption_present(decision):
        draft.add(mark_assumption(decision.assumptions))
    if decision_requires_options(decision):
        draft = package_as_options(draft)
    if draft_too_abstract(draft):
        draft.add(example_for_context(decision))
    if audience_needs_plain_language(decision):
        draft.add(plain_language_sentence(draft))
    if user_needs_actionable_output(decision):
        draft.add(next_step(decision))
    return draft
```

This is the minimum runtime shape of **Generate**. A real implementation may simplify or expand individual generation functions, but the stage should still produce a first draft from the plan, shape it for the intended audience and task, make key assumptions visible where needed, and express uncertainty in a form that remains usable.

Why Generate Matters

A well-run **Generate** stage keeps AVA from confusing clarity with display. It gives the runtime a draft that is already trying to be usable before validation begins.

Later stages work best when they are correcting a strong first attempt rather than rescuing an incoherent or overproduced one.

Generate is where the plan becomes language, but it's not where the system earns the right to stop thinking. The draft still has to survive validation. What this stage contributes is disciplined first expression: a response that already knows what it's trying to do.

Human Parallel

A good human answer usually tries to say the thing clearly the first time. It may still need revision, but it does not begin by stalling, performing intelligence, or circling the point.

Generate is that first committed attempt.

6. Validate

Validate is the enforcement stage of the runtime.

It checks the draft against the behavioral contract established earlier in the loop and determines whether the response is fit to release, needs revision, should narrow, or must stop short entirely.

This stage exists because fluent drafting is not a sufficient guarantee of good behavior. A response can be readable and still be unsafe, unsupported, structurally thin, emotionally distorted, prematurely synthetic, repetitive, or unable to stop.

Validate is what prevents AVA from mistaking a plausible draft for a finished answer.

At minimum, **Validate** should produce a corrected draft and a release decision.

That means the stage must do more than score the response—it must be able to intervene. If the draft is overclaiming, drifting, escalating, looping, or speaking with more certainty than it has earned, **Validate** must be able to trim, revise, downshift, block, or redirect it before release.

Core Validator Sequence

The ordered validator sequence remains:

Containment → **Drift & Layer Balance** → **Horizon Arcs** → **Recursion Control** → **Language Hygiene** → **Closure**

That order is binding. Safety, scope, and grounding failures must be caught before style cleanup. Progression must be checked before the answer is allowed to conclude. **Closure** is last because the system should stop only after the rest of the draft has been brought into proportion.

Modules and Frameworks

The **Validate** stage is primarily shaped by the following components:

- **Containment** (p. 14), which checks whether the draft is safe, scoped, and sufficiently grounded for the claim it is making.

- **Drift & Layer Balance** (p. 14), which checks whether the response has wandered off-task or become dominated by **Performance**, **Emotion**, or **Structure**.
- **Horizon Arcs** (p. 14), which checks whether the draft has jumped ahead into recognition, continuity, or synthesis before earning those moves.
- **Recursion Control** (p. 14), which prevents repetition, unnecessary continuation, and self-reinforcing loops.
- **Language Hygiene** (p. 14), which removes apology spirals, filler, canned phrasing, and language that revisits earlier material without adding meaning.
- **Closure** (p. 14), which checks whether the response has actually arrived and can be allowed to end.
- **Layer Analysis and Rebalancing** (p. 24), which estimates underweight and overweight layers and applies corrective adjustments.
- **Horizon Accounting and Gate Memory** (p. 26), which uses prior state to determine whether advanced moves are actually supported across turns.

Validation Extensions

These extensions are optional, but often useful in tighter-control systems:

- **Numeracy and Units Check** (p. 36), which verifies arithmetic, conversions, and units before the answer is released.
- **Contradiction Check** (p. 36), which compares the draft against itself, the evidence package, and prior state for internal mismatch.
- **Confidence Calibration** (p. 37), which adjusts certainty level to match what is actually supported.
- **Refusal Quality Check** (p. 37), which improves refusal behavior so it is bounded, honest, and useful rather than empty or theatrical.

What Validate Should Produce

Before AVA moves into **Close**, the **Validate** stage should be able to answer the following questions:

1. Is the draft safe and appropriately scoped?

2. Is it still on the user's actual brief?
3. Are **Performance**, **Emotion**, and **Structure** in usable proportion?
4. Has the draft moved too quickly into synthesis or continuity?
5. Is it repeating itself or continuing without new value?
6. Is the language cleaner than the raw generation pass?
7. Is certainty matched to support?
8. Does the draft now deserve release?

The point is not to make the answer sterile, but fit for use.

Illustrative Pseudocode

The functions shown here represent enforcement and correction steps rather than fixed implementations.

In practice, they may be realized through rule-based checks, scoring functions, model-assisted revision passes, or hybrids that combine these approaches. What matters is that **Validate** does not merely observe the draft—it must be able to modify it.

This is the stage where the grammar becomes enforceable rather than advisory.

```

Validate(draft, decision, context, evidence, state):
  if unsupported_or_high_risk(draft, evidence):
    draft = downshift_or_constrain(draft)
  if off_brief(draft, decision):
    draft = reanchor_to_task(draft)
  layer_mix = inspect_layers(draft)
  draft = adjust_layer_mix(draft, layer_mix)
  horizon_shape = inspect_horizon_progression(draft)
  if horizon_jump_detected(horizon_shape):
    draft = downshift_horizon(draft)
  if advanced_move_unearned(draft, state):
    draft = reduce_advanced_claims(draft)
  if looping_without_progress(draft):
    draft = compress_or_stop(draft)
  draft = strip_filler_and_templates(draft)
  if numeric_claims_present(draft):
    draft = verify_numbers_and_units(draft)
  if contradiction_found(draft, evidence, state):
    draft = resolve_or_mark_contradiction(draft)
  draft = calibrate_confidence(draft, evidence)
  if refusal_needed(draft):

```

```
draft = improve_refusal_quality(draft)
ready_to_close = answer_has_arrived(draft, decision)
return ValidatedDraft(
    draft=draft,
    ready_to_close=ready_to_close
)
```

This is the minimum runtime shape of **Validate**. A real implementation may simplify or expand individual checks, but the stage should still be able to constrain risk, restore proportion, enforce progression limits, reduce repetition, calibrate certainty, and determine whether the draft is ready to close.

Why Validate Matters

A well-run **Validate** stage is what makes AVA a runtime system rather than a writing preference.

It's the point where the framework governs behavior instead of merely describing it. Without validation, every earlier stage remains advisory. With validation, the system gains enforcement.

Many conversational failures are not failures of intelligence. They are failures of stopping, pacing, proportion, confidence, and grounded release. **Validate** is where those failures become visible while there is still time to correct them.

Human Parallel

A good human answer often includes an internal second pass before it is said out loud: does this actually hold up, am I overstating it, am I still on the point, have I already answered the question? **Validate** is that second pass made structural.

7. Close

Close is the final active stage of the runtime. It determines whether the validated draft has actually completed the work of the turn and whether the response can now be released without further continuation.

This is the stage where AVA treats stopping as part of correct behavior rather than as the accidental absence of more words.

This stage exists because many conversational systems know how to begin and continue, but not how to finish. Even after the answer has arrived, they keep extending the exchange with redundant explanation, softened restatement, optional next steps that feel compulsory, or speculative continuation that adds little beyond motion.

Close prevents that drift at the endpoint. It asks a simple question: *is the work done?*

At minimum, **Close** should produce a final release decision.

That means the stage must determine whether the validated draft has landed at the right size, with the right degree of completion, and whether any final handoff is genuinely useful rather than merely habitual.

Modules and Frameworks

The **Close** stage is primarily shaped by the following components:

- **Closure** (p. 14), which confirms that the answer has reached a natural endpoint and should be allowed to stop there.
- **Response Surface Rules** (p. 18), which apply the final check on size, density, and directness so the response ends without decorative spill.
- **Stopping Rule** (p. 37), which defines what counts as enough for the current task so the response does not continue by momentum alone.
- **One-Step-Now** (p. 37), which adds one low-friction next move when a final handoff would make the response more useful.
- **Time-Frame Sorter** (p. 37), which organizes the ending into now, next, and later when that improves usability at the moment of release.

What Close Should Produce

Before AVA releases the answer, the **Close** stage should be able to answer the following questions:

1. Has the actual task of the turn been completed?
2. Is there any meaningful value in continuing?
3. Would a final next step help, or would it only prolong the exchange?
4. Does the ending arrive cleanly, or does it soften into unnecessary extra language?
5. Is the response ending at the right size for the request?

The point is clean completion.

Illustrative Pseudocode

The functions shown here represent release and stopping decisions rather than fixed implementations.

In practice, they may be realized through rule-based stop conditions, task-completion checks, model-assisted judgments about usefulness, or hybrids that combine these approaches.

What matters is that **Close** does not treat release as automatic. It must determine whether the work has been completed, whether any final handoff is genuinely useful, and whether the response can end without further continuation.

```
Close(validated_draft, decision, context):
if task_completed(validated_draft, decision):
    ready_to_release = True
else:
    ready_to_release = False
if answer_has_arrived(validated_draft, decision):
    validated_draft = mark_closed(validated_draft)
if user_needs_action_handoff(decision):
    validated_draft = append_one_step_now(validated_draft)
if closing_needs_temporal_handoff(decision):
    validated_draft = add_time_frame_handoff(validated_draft)
validated_draft = enforce_final_surface_rules(validated_draft)
return FinalResponse(
    draft=validated_draft,
    released=ready_to_release
)
```

This is the minimum runtime shape of **Close**. A real implementation may simplify or expand individual release checks, but the stage should still determine whether the work of the turn is complete, whether any handoff improves usability, and whether the response can end cleanly at the right point.

Why Close Matters

A well-run **Close** stage protects the user from invisible labor at the endpoint of the exchange.

Without it, the user is left to decide where the answer should have stopped, which extra paragraph mattered, whether the optional next step was actually necessary, and whether the system is finished or merely pausing before more continuation. **Close** removes that burden by making completion explicit.

The end of a response affects trust as much as the beginning. A system that consistently knows when to stop feels more reliable, more usable, and less extractive of attention. It feels like a tool rather than a conversation machine trying to remain in motion.

Human Parallel

A good human answer often ends when the point has landed. It does not keep speaking just to prove it could have said more. **Close** is the part of AVA that makes that restraint structural.

8. State Writeback

State Writeback is the final system-level step of the runtime.

It updates conversation state after the response has been completed so the next turn begins with the right continuity rather than accumulated noise. If **Sense** reads the prior position of the exchange, **State Writeback** determines what that position becomes.

This stage exists because continuity should not depend on dragging the full transcript forward forever. A system that carries everything carries too much; a system that carries nothing keeps starting over.

State Writeback preserves the minimum information needed for the next correct move: what was established, what remains unresolved, what was grounded, what defaults are now active, and what validator or horizon conditions should remain visible in later turns.

At minimum, **State Writeback** should produce a compact next-state object. That object should preserve position, not prose. Its job is not to remember everything that was said. Its job is to remember what the runtime now knows, what it still owes, and what should shape the next turn.

Modules and Frameworks

The **State Writeback** stage is primarily shaped by the following components:

- **State Tracking** (p. 20), which writes forward only the variables needed to preserve continuity across turns.
- **Horizon Accounting and Gate Memory** (p. 26), which carries forward what level of progression has actually been earned so later turns do not restart too high.
- **Motif Compression / Small Recognizers** (p. 45), which preserve recurring patterns in compressed form rather than hauling the raw exchange forward.

What State Writeback Should Produce

Before the system begins another turn, **State Writeback** should be able to answer the following questions:

1. What has now been established strongly enough to carry forward?

2. What remains unresolved?
3. What horizon movement has been earned?
4. What validator interventions still matter?
5. What user-level defaults should persist?
6. What recurring motifs should remain visible?
7. What should be dropped rather than carried forward?

The point is not memory for its own sake, it's useful continuity with minimal waste.

State Writeback Structures

State Writeback doesn't carry forward the full transcript—it writes a compact representation of what now matters for the next turn. The following structures are typical ways to organize that state.

They're local to this stage: examples of how continuity can be preserved without hauling language forward.

Confirmed Grounding Register stores what has been grounded strongly enough to be treated as established inside the thread. This gives later turns a stable factual floor without forcing the runtime to re-derive the same supported material.

Open Tensions Register preserves unresolved tensions, tradeoffs, or active questions that still matter. These are often the pieces that shape the next turn even when the current response has otherwise completed its work.

Validator Event Log stores recent validator interventions that may still matter, such as repeated drift, unsupported claims, or recursion pressure. This allows later turns to respond to short-term behavioral patterns rather than treating each new message as though it arrived in a vacuum.

User Defaults and Preferences carries forward durable local preferences such as size, depth, stance, or formatting defaults when those have become part of the runtime. That reduces the need to re-infer the same preferences repeatedly across a thread.

Illustrative Pseudocode

The functions shown here represent writeback and compression steps rather than fixed implementations.

In practice, they may be realized through structured state objects, rule-based persistence logic, summarization layers, or hybrids that combine these approaches. What matters is that **State Writeback** carries forward only what improves the next turn.

It should preserve grounded facts, live tensions, earned progression, durable defaults, and recent validator signals, while dropping transcript-heavy residue that does not improve—and may degrade—future behavior.

```
next_state = {}

# carry forward current mode and durable user defaults
next_state["mode"] = context.mode
next_state["defaults"] = update_user_defaults(state, context, final_response)

# preserve what is now established strongly enough to stand on
next_state["confirmed_grounding_register"] = collect_confirmed_grounding(
    evidence, final_response
)

# preserve unresolved questions, tensions, or tradeoffs
next_state["open_tensions_register"] = collect_unresolved_tensions(
    final_response, state
)

# compress recurring motifs rather than carrying transcript bulk
next_state["motifs"] = compress_recurring_patterns(
    context, final_response, state
)

# preserve earned progression limits
next_state["horizon_floor"] = compute_horizon_floor(final_response, state)
next_state["horizon_ceiling"] = compute_horizon_ceiling(final_response,
    state)

# preserve validator interventions that may still matter next turn
next_state["validator_event_log"] = collect_validator_events(
    validation_result, state
)

# drop anything that does not improve the next turn
next_state = prune_nonessential_state(next_state)

return next_state
```

This is the minimum runtime shape of **State Writeback**. A real implementation may simplify or expand individual persistence functions, but the stage should still preserve only the information that improves the next turn and discard the rest.

Memory Efficiency and Infrastructure Note

This stage is also where some of the AVA Framework's efficiency gains may appear in real deployments.

If the runtime carries forward compact state, motifs, constraints, and decisions instead of repeatedly hauling raw transcript forward, the active context footprint can remain smaller over long sessions.

In some implementations, that can reduce context-window pressure, slow KV-cache growth, and lower memory load on the serving stack. These are not automatic properties of the grammar by themselves; they are testable infrastructure effects of a system that writes back only what matters and lets the rest go.

Why State Writeback Matters

A well-run **State Writeback** stage is what makes AVA feel continuous without becoming bloated. It lets later turns begin from a real position rather than from total amnesia or transcript hoarding.

Long threads do not usually fail all at once. They fail by accumulation: too much context, too much residue, too much repeated material, and too little distinction between what should be carried forward and what should have been dropped.

State Writeback is the stage that resists that accumulation. It gives the next turn a floor, a few live tensions, some earned progression, and a compact memory of what the runtime now knows. Nothing more is required.

Human Parallel

A good conversation does not remember every sentence. It remembers what was established, what still matters, and what the next move has to account for. **State Writeback** is that discipline made structural.

Part III — Integration Profiles

Part III translates the AVA Framework from general runtime architecture into deployment shape.

Part I defines the system conceptually.

Part II shows the system in motion as a blueprint.

This part answers a different question: how the same conversational grammar should be configured when it is implemented in environments with different stakes, tolerances, and operating conditions.

These profiles don't create separate frameworks—AVA remains one core framework. What changes across profiles is the tuning: which modules are necessary, which thresholds become stricter, which defaults should govern the runtime, and which kinds of failure carry the highest cost in that environment.

The runtime should not be configured identically in every setting.

A social or entertainment system may tolerate more looseness, play, and stylistic variance. A tutoring system may require clearer pacing, stronger clarification behavior, and more visible progression. A clinical, legal, or financial system may require earlier containment, stricter grounding, narrower claims, and stronger confidence calibration. A machine-to-machine integration may suppress most human-facing performance features while preserving the same **Planner Loop**, evidence logic, and validation order.

These integration profiles are deployment postures, not personality presets.

Each profile shows how the same runtime should be adjusted to match a different environment without losing the AVA contract.

How to Read the Profiles

Each profile follows the same structure:

- a short description of the environment
- the primary runtime priorities in that setting
- the modules or recognizers most likely to matter
- the failure modes that deserve the most attention

- the practical posture AVA should take there

The goal is to make adaptation legible. A team should be able to look at a profile and quickly see which parts of the blueprint are essential, which are optional, and where stricter controls or looser defaults make sense.

Profiles Included in Part III

1. **Consumer / Social / Entertainment** allows higher stylistic flexibility, greater tolerance for play and looseness, and continued emphasis on drift control and closure.
2. **Enterprise / Internal Tools** prioritizes stronger scope control, traceability, operational consistency, and lower tolerance for ambiguity or wasted motion.
3. **Tutoring / Coaching / Education** emphasizes guided progression, clarification, pedagogical pacing, and stronger support for examples and sequencing.
4. **Clinical / Legal / Financial** requires stricter grounding thresholds, narrower claims, earlier containment, and stronger confidence calibration.
5. **Machine-to-Machine / System Integrations** minimizes the human-facing performance layer while preserving strong evidence, sequence, and validation logic, with output optimized for execution rather than conversation.

What Changes Across Profiles

The underlying runtime remains. Every profile still preserves:

- ordered execution
- grounding when required
- active validation
- proportion maintenance
- explicit closure

What changes is the configuration around those commitments:

- the default response surface
- the strictness of containment

- the aggressiveness of retrieval and evidence gates
- the relative importance of examples, action scaffolds, or recognizers
- the cost of false confidence, drift, or premature synthesis

A useful way to understand the profiles is this: the runtime remains constant, but the tolerances move.

Why Integration Profiles Matter

A framework that cannot adapt to different environments remains theoretical.

A framework that adapts too freely stops being itself.

The integration profiles hold the middle ground. They show how AVA can be tuned to real products and real constraints without dissolving into vague best practices or collapsing into one narrow use case.

That is what makes deployment legible: not only knowing what the system is, but knowing what it should protect against, what it should optimize for, and how it should behave in a specific setting.

1. Consumer / Social / Entertainment

This profile covers public-facing systems used in ordinary, lower-stakes interaction, where the exchange may be casual, exploratory, playful, or expressive.

The runtime does not become less structured in this environment; it becomes lighter at the surface. The system may allow more stylistic range, more personality in expression, and more tolerance for imaginative framing, but it should still preserve the same underlying order of interpretation, planning, grounding, validation, and closure.

The main objective in this profile is to keep the interaction lively without letting liveliness replace discipline.

A consumer-facing system can tolerate more looseness than an enterprise or higher-stakes deployment, but it still needs protection against drift, overperformance, false confidence, emotional over-identification, and endless continuation. The answer should feel readable, responsive, and emotionally legible, while still landing clearly and stopping once the point has been made.

Best Fit

The most important components in this profile are the ones that preserve proportion while keeping the surface easy to use. The system needs enough structure to resist drift and enough flexibility to avoid feeling stiff or overcontrolled.

In practice, that means prioritizing modules that keep answers direct, readable, balanced, and capable of stopping once the point has landed.

Best fit: Layer Analysis and Rebalancing — Response Surface Rules — Answer-First Shaper — Example Injector — Plain-Language Pass — Mode Setter — Four Levers — Motif Spotting / Small Recognizers

Good Fit

These components become useful when the interaction needs a little more structure than casual exchange alone provides.

They help when the user is choosing among options, needs a practical next move, or would benefit from clearer sequencing without turning the answer into a rigid task flow.

Good fit: Option Packager — Action Step Finisher — One-Step-Now — Three Horizons — Signal → Story → Scar

Possible / More Selective

These components matter when the interaction starts to lean toward factual exposure, time sensitivity, stronger refusal needs, or claims that carry more consequence than the default consumer setting usually assumes.

They should be available, but not treated as part of the default surface behavior in every exchange.

Possible: Explicit Grounding Triggers — Freshness Gate — Confidence Calibration — Refusal Quality Check

Practical Posture

In this profile, AVA should feel flexible at the surface and disciplined underneath.

It can allow more range in tone, humor, metaphor, or stylistic expression than other profiles, but it should still regulate the same deeper risks. The answer should remain proportionate, readable, and capable of stopping once the point has landed.

The design challenge here is to let the system feel alive without letting it become slippery. That means preserving strong drift control, clean closure, and enough layer balance that **Performance** does not consume the whole interaction.

Human Parallel

This version of AVA should be pleasant to use without becoming all vibe and no floor.

2. Enterprise / Internal Tools

This profile covers systems used inside organizations for work that is operational, collaborative, or decision-supportive rather than purely conversational.

These environments often involve recurring tasks, internal documents, process handoffs, team coordination, and a lower tolerance for wasted motion. The interaction doesn't need to feel cold, but it does need to feel reliable.

The key difference in this profile is that scope control, traceability, and consistency take priority over expressive range. The runtime should remain readable and humane, but it should bias toward clarity, boundedness, and execution. A strong enterprise deployment should make it easier for users to complete work, verify what the system is doing, and reuse outputs without having to re-stabilize the interaction every time.

The main objective in this profile is to reduce invisible labor.

Enterprise users should not have to keep re-asking, re-scoping, rechecking assumptions, or cleaning up sprawling answers before they can act. The system should surface constraints early, preserve task clarity across turns, and return outputs that are easy to verify, circulate, and use.

In this setting, a helpful answer is not enough; the answer should also be stable under reuse.

Best Fit

The most important components in this profile are the ones that turn interpretation into execution. The system needs to clarify gaps before moving, decompose work into a usable plan, expose assumptions and constraints early, and keep answers bounded enough to support recurring workflows.

It also needs strong checks on evidence, contradiction, confidence, and stopping so the output remains reliable when it is reused downstream.

Best fit: Clarify-or-Answer Planner — Query-Plan Decomposer — Assumption Ledger — Missing Constraint Finder — Five Switches — Three Horizons — Time-Frame Sorter — Retrieval Gate / Context Diet — Evidence Sufficiency Check — Contradiction Check — Confidence Calibration — Stopping Rule — One-Step-Now — State Tracking

Good Fit

These components become especially useful when the work depends on consistent formatting, internal source control, or operational precision across different kinds of tasks. They help the system settle into the right mode quickly, package outputs more cleanly, and catch numerical or freshness-related issues that would otherwise create avoidable rework.

Good fit: Mode Setter — Four Levers — Action Step Finisher — Option Packager — Canon Resolver — Freshness Gate — Numeracy and Units Check

Possible / More Selective

These components matter in narrower cases where the internal workflow becomes more explanatory, more humanly charged, or more audience-sensitive than the default enterprise setting usually assumes. They can improve clarity or refusal quality when needed, but they should not dominate the default posture of the system.

Possible: Signal → Story → Scar — Example Injector — Plain-Language Pass — Refusal Quality Check

Practical Posture

In this profile, AVA should feel worklike rather than theatrical.

The answer should usually arrive in a bounded, usable form with visible assumptions, explicit next moves, and as little interpretive burden as possible. Retrieval and contradiction checks carry more weight here because enterprise users often need outputs they can circulate or build on without manually rechecking every sentence.

The strongest version of AVA in this setting isn't the most impressive-sounding one, but the one that reduces invisible labor: less re-asking, less re-scoping, less transcript babysitting, and less cleanup after the answer arrives.

Human Parallel

This version of AVA should feel like a sober operator: clear, steady, on-task, and finished when the work is done.

3. Tutoring / Coaching / Education

This profile covers systems used for learning, teaching, guided explanation, practice, and developmental support.

The common requirement across tutoring, coaching, and education is not just answer delivery, but movement through understanding. The system is being used to help someone learn, clarify, practice, or progress, whether the domain is academic, personal, or professional.

The key difference in this profile is that progression carries as much weight as correctness.

A strong tutoring or coaching system should not compress every problem into the shortest possible answer, nor should it flood the user with explanation that outruns readiness. The runtime should preserve enough structure to keep the exchange grounded while also supporting staged understanding, clarification, and one-step-at-a-time movement.

The main objective in this profile is to match the response to the learner's or client's current position. The system should make reasoning visible without overwhelming the user, use examples and sequencing well, and help the person move forward with durable footing rather than borrowed fluency.

In this setting, the risk is not only getting the answer wrong. It's answering too fast, synthesizing too early, or carrying the user to closure without helping them build their own understanding.

Best Fit

The most important components in this profile are the ones that regulate pace, preserve progression, and keep the next step usable.

The system needs to identify what kind of learning exchange is taking place, decide whether clarification is needed, break the task into manageable structure, and keep the response within the learner's or client's next usable horizon. It should also make explanation concrete enough to follow without turning every exchange into a lecture.

Best fit: Mode Setter — Clarify-or-Answer Planner — Query-Plan Decomposer — Three Horizons — Example Injector — Plain-Language Pass — Answer-First Shaper — Action Step Finisher — One-Step-Now — Horizon Progression — Horizon Accounting and Gate Memory — State Tracking

Good Fit

These components become especially useful when the learning or coaching process depends on clearer execution shape, better pacing across phases, or more explicit adjustment to what the user is missing.

They help the system notice blocked movement, rebalance the response when it becomes too thin or too dense, and package alternatives more clearly when several paths are available.

Good fit: Five Switches — Time-Frame Sorter — Missing Constraint Finder — Layer Analysis and Rebalancing — Option Packager — Four Levers

Possible / More Selective

These components matter in narrower cases where the exchange becomes more emotionally layered, more causally complex, or more dependent on external support than the default tutoring or coaching setting usually assumes.

They can improve depth and calibration when needed, but they should not become the default center of gravity for every learning interaction.

Possible: Signal → Story → Scar — Layered Cause — Retrieval Gate / Context Diet — Evidence Sufficiency Check — Confidence Calibration

Practical Posture

In this profile, AVA should act as a paced guide rather than a pure answer engine. It should still be capable of direct answers when that is what the user needs, but it should also know when to slow down, clarify the frame, add a concrete example, break the task into steps, or keep the response within the next usable horizon.

The strongest tutoring or coaching deployment is one that makes progress feel reachable without pretending every exchange must become a lecture. The answer should help the user move, not just admire the map.

Human Parallel

This version of AVA should feel like a good teacher or coach: clear, patient, paced, and interested in whether the user can actually take the next step.

4. Clinical / Legal / Financial

This profile covers higher-stakes environments where the cost of unsupported language is materially higher.

These settings differ from general assistance not because the runtime changes in kind, but because the tolerances narrow. Claims must be better grounded, uncertainty must appear earlier, refusal quality carries more weight, and the system should be much less willing to improvise past missing support.

The key difference in this profile is that AVA should bias toward containment before conversational smoothness.

A system in these domains can still be clear, humane, and readable, but it should not preserve flow at the expense of support, scope, or appropriate caution. The right answer in these settings is often narrower than the most conversationally satisfying one.

The main objective in this profile is to keep the system inside the evidentiary and practical floor of the task. It should surface uncertainty, constraints, and missing support early; prefer narrower claims over broader but shakier ones; and make refusal, redirection, and limitation behavior useful rather than abrupt.

In this setting, the risk is not only being wrong—it's sounding authoritative without having earned that authority.

Best Fit

The most important components in this profile are the ones that prevent unsupported language from hardening into advice, interpretation, or implied certainty.

The system needs strong grounding discipline, clear trigger logic for when support becomes mandatory, and reliable checks on sufficiency, freshness, contradiction, and confidence. It also needs refusal and stopping behavior that remains specific and usable when the answer has to narrow or stop short.

Best fit: Grounding Behavior — Explicit Grounding Triggers — Evidence Sufficiency Check — Evidence Gate — Freshness Gate — Canon Resolver — Contradiction Check — Confidence Calibration — Refusal Quality Check — Missing Constraint Finder — Stopping Rule

Good Fit

These components become especially useful when the task still requires structure, sequencing, or state awareness in addition to stronger containment.

They help the system clarify what is being asked, organize the answer more carefully, track ongoing constraints across turns, and catch numerical or temporal details that could materially change the reliability of the response.

Good fit: Clarify-or-Answer Planner — Query-Plan Decomposer — Numeracy and Units Check — Time-Frame Sorter — State Tracking — Layer Analysis and Rebalancing

Possible / More Selective

These components matter in narrower cases where clarity, accessibility, or next-step guidance can improve usability without loosening the evidentiary floor.

They should be applied carefully, because in these environments warmth, simplification, or action framing can become misleading if they imply more confidence or completeness than the system has earned.

Possible: Plain-Language Pass — Example Injector — Action Step Finisher — One-Step-Now — Four Levers — Signal → Story → Scar

Practical Posture

In this profile, AVA should feel careful without becoming evasive.

It should move early toward support, qualification, narrowing, or refusal when the floor is missing. A response may still be readable and humane, but those qualities should not lead the system into implying more confidence, authority, or completeness than it has earned.

The strongest version of AVA in these settings helps the user understand what can be said, what cannot yet be said, what remains uncertain, and what kind of next move is actually appropriate. That may include redirection toward a human professional, a narrower explanation of possibilities, or a refusal that is specific enough to be useful.

Human Parallel

This version of AVA should feel like a careful professional: clear, bounded, properly cautious, and unwilling to pretend certainty where certainty has not been earned.

5. Machine-to-Machine / System Integrations

This profile covers environments where AVA is not primarily speaking to a human reader, but operating between systems, services, tools, or structured downstream processes.

In these contexts, the conversational grammar still matters, but its expression changes. The runtime remains ordered, grounded, and validated, while most human-facing stylistic features become secondary or disappear entirely.

The key difference in this profile is that the system should optimize for execution shape rather than conversational feel.

The output must still be coherent, but coherence here means something narrower and more operational: structured inputs are interpreted correctly, plans are formed cleanly, evidence is checked when needed, outputs remain bounded, and the system stops without producing narrative spill that downstream tools did not ask for.

The main objective in this profile is to preserve runtime discipline while making outputs easier to parse, verify, and route. The system should minimize ambiguity in what it emits, keep claims and state transitions explicit, enforce grounding and contradiction checks where required, and suppress unnecessary tone, flourish, or elaboration.

In this setting, the risk is not only saying something incorrect. It's producing output that forces downstream interpretation, propagates unsupported claims into later actions, or continues after the required work has already been done.

Best Fit

The most important components in this profile are the ones that preserve strict order, bounded state, evidentiary discipline, and clean stopping under structured execution.

The system needs to interpret inputs carefully, choose a response path deliberately, retrieve only what is required, validate what it emits against support and state, and stop once the required output has been produced. These are the components that make AVA behave like reliable runtime infrastructure rather than a conversational surface.

Best fit: Runtime Contract — Grounding Behavior — Explicit Grounding Triggers — Retrieval Gate / Context Diet — Evidence Sufficiency Check — Evidence Gate — Contradiction Check — Numeracy and Units Check — Stopping Rule — State Tracking — Horizon Accounting and Gate Memory

Good Fit

These components become especially useful when the integration has to resolve ambiguity, manage structured sequencing, or respect source authority across multiple systems. They help the runtime clarify when it has enough to proceed, decompose work into cleaner internal steps, detect missing constraints before action is taken, and maintain better control over canon, freshness, and execution order.

Good fit: Clarify-or-Answer Planner — Query-Plan Decomposer — Canon Resolver — Freshness Gate — Missing Constraint Finder — Time-Frame Sorter — Five Switches

Possible / More Selective

These components matter in narrower cases where a human may still inspect the output, where readability has some downstream value, or where recurring patterns benefit from lighter recognizer support. They should be used selectively, because in this profile the default priority is machine readability and bounded execution, not conversational richness.

Possible: Response Surface Rules — One-Step-Now — Option Packager — Motif Spotting / Small Recognizers

Practical Posture

In this profile, AVA should feel quiet, exact, and bounded.

The system does not need to sound vivid or relational unless a human is explicitly in the loop. It needs to preserve the same internal discipline while making outputs easier to parse, verify, and route. That usually means shorter emission, stronger structure, narrower claims, and clearer state handling.

The strongest version of AVA in this setting behaves like a reliable runtime intermediary. What remains from the broader grammar is not the surface style, but the order of operations, the grounding discipline, the enforcement layer, and the stopping behavior. Those are the parts that matter most when machines are speaking to machines.

Human Parallel

This version of AVA should feel like clean infrastructure: precise, compact, and dependable, with no extra conversation leaking into the wire.

Closing Note on Integration Profiles

These profiles show how the same conversational grammar can be tuned for different environments without losing its core runtime contract—the loop remains the same.

What changes is the posture: which risks carry the highest cost, which modules are worth keeping, which thresholds should tighten, and what kind of behavior the environment can tolerate.

That's also the point of the document as a whole. AVA is not presented here as a finished product to be copied exactly. It's a blueprint that can be implemented, tested, modified, narrowed, expanded, or partially adopted in real systems.

One team may begin with the core **Planner Loop** and a small validator layer, then find immediate value. Another may implement the full blueprint, tune it to a specific profile, and instrument the runtime in production. The standard is whether the system becomes more proportionate, more grounded, less prone to drift, and easier to use in practice.

That's why the framework is modular, why the blueprint is shown in full, and why the hypotheses in the next part matter. The right response to AVA is not belief, but testing.

Use the profile that fits most closely, modify what doesn't, delete what you do not need, keep what improves the runtime, create new modules where the environment calls for them.

If the grammar works, that should become visible in the behavior of the system itself. If it doesn't, that should become clear as well.

Part IV — Hypotheses for Evaluation

Part IV defines how the AVA Framework should be evaluated.

AVA operates at the interaction layer. Its claims are therefore behavioral: that a system can stay on task more reliably, ground what must be grounded, reduce drift, reach cleaner completion, and require less manual correction from the user. Those are runtime behaviors, and they should be measurable.

The purpose of this part is to state what AVA predicts, how those predictions can be tested, what should be measured, and what kinds of results would count for or against the framework. This includes both formal evaluation methods and lighter assessment surfaces that make conversational behavior easier to inspect in practice.

Part IV is written for builders, researchers, and teams deciding whether AVA improves a real system. It does four things:

1. It defines the primary and secondary hypotheses of the framework.
2. It outlines practical evaluation design.
3. It identifies useful measures for runtime behavior.
4. It introduces the **Coherence Receipt** as a supplementary way of making conversational quality visible in a fast, readable form.

The evaluation posture in this part is comparative.

The useful question isn't just whether a single AVA response sounds good in isolation but whether a system running AVA behaves better than an appropriate baseline on tasks that matter in practice. In most cases, that means comparing the same model without AVA, the same model with a lighter AVA layer, the same model with fuller AVA integration, or AVA-shaped behavior against the team's current orchestration or policy layer.

This part also assumes that evaluation should include the kinds of tasks most likely to expose runtime weakness. Long threads, ambiguous asks, emotionally charged prompts, planning tasks, document-bound interpretation, time-sensitive questions, and situations with incomplete grounding are more useful than easy cases that flatter the system. If AVA improves the interaction layer, that improvement should become visible under pressure rather than only in ideal conditions.

The framework is modular, so the evaluation can be modular as well.

Some teams may evaluate the full runtime. Others may test the **Planner Loop**, the **Validator Suite**, selected modules, recognizers, or prompt-layer AVA behavior against fuller implementations.

Some may care most about efficiency and closure. Others may care most about grounding, drift resistance, actionability, or cross-turn continuity.

This part is written to support both kinds of use: evaluation of AVA as a whole and evaluation of the parts that can be adopted independently.

That is also where the **Coherence Receipt** becomes useful. Benchmarks and task metrics show whether performance changes. The receipt provides a lighter assessment surface for how that change appears in the exchange itself: whether a response held its shape, where it slipped, and which parts of the runtime were carrying the interaction well or failing under pressure.

It's not a substitute for formal measurement, but it is a practical instrument for making conversational behavior more visible during comparison, review, and iteration.

The sections that follow move from evaluation posture to hypotheses, from hypotheses to test design, and from test design to measurement. They then define what success, mixed results, and failure should look like, including cases where only part of the framework proves useful.

The aim of this part is straightforward: to make AVA testable enough that teams can determine whether it improves behavior in use.

1. Evaluation Posture

AVA should be evaluated comparatively, not in isolation.

What matters is how a system behaves with AVA in place compared to a relevant baseline under real conditions. A single response, even if it reads well, does not say much on its own. The useful signal comes from comparison.

In most cases, that comparison should stay controlled:

- the same model without AVA
- the same model with a lighter prompt-layer version
- the same model with fuller AVA integration
- or AVA-shaped behavior against the current orchestration or policy layer

The goal is to isolate what the runtime is changing, rather than introducing unrelated variables.

The evaluation should also focus on the kinds of tasks where interaction quality tends to break down. Long threads, ambiguous requests, emotionally charged prompts, planning tasks, document-bound interpretation, time-sensitive questions, and situations with incomplete or unstable information are all more revealing than clean, well-formed prompts. These are the conditions where drift, overextension, unsupported claims, and poor stopping behavior usually show up.

Testing only in ideal conditions will not tell you much. The framework needs to be observed under pressure, where the interaction layer is actually doing work.

A system that performs well on simple prompts but requires constant correction in these conditions is not stable. A system that holds shape, stays grounded, and reaches usable completion with less intervention is.

That difference is what this part is designed to surface.

2. Primary Hypotheses

The framework makes four primary claims. These are the first things to test and the most important outcomes to track. If AVA is improving the interaction layer, the effect should show up here.

2.1 Efficiency

Hypothesis: AVA improves conversational efficiency by reducing unnecessary continuation, repeated clarification loops, output padding, and correction labor for the user.

Test this by: comparing how much interaction is required to reach a usable result across matched tasks. Keep the comparison controlled: the same model without AVA, the same model with a lighter AVA layer, or the same model with fuller AVA integration.

Look for: fewer output tokens per resolved task, fewer turns to task completion, fewer user correction turns, less repeated re-scoping, and fewer cases where the user has to ask the system to summarize, tighten, or stop.

In longer sessions: look for lower transcript residue, less repeated cleanup work, and less accumulation of unnecessary continuation across turns.

What would count as support: the AVA-shaped system reaches usable completion with less wasted interaction while preserving answer quality.

What would count against it: the system becomes more ceremonial, more expensive, or more time-consuming without reducing user effort or improving completion.

2.2 Grounding

Hypothesis: AVA improves grounding discipline by making unsupported claims less likely to pass through as fluent authority.

Test this by: running comparisons on tasks that depend on external facts, document interpretation, named source material, current or unstable information, or higher-stakes explanation. Include cases where the evidence is present, incomplete, contradictory, or missing.

Look for: earlier uncertainty marking, better support-seeking behavior, fewer unsupported factual claims, narrower answers when support is missing, and fewer confident misreads of provided material.

Where it should show up most clearly: document-bound tasks, factual prompts with hidden evidence gaps, time-sensitive questions, and situations where the model is tempted to keep speaking past what it actually knows.

What would count as support: the system draws a cleaner boundary between what is supported, what is inferred, and what remains unknown, and it behaves accordingly.

What would count against it: unsupported claims continue to pass through in polished form, uncertainty appears late or not at all, or retrieval overhead increases without improving factual discipline.

2.3 Drift

Hypothesis: AVA reduces conversational drift by keeping the system closer to the actual task and reducing side-quests, repetition, overextension, and topic slippage across long or dense exchanges.

Test this by: comparing AVA-shaped and baseline systems on prompts that invite overreach: long threads, ambiguous requests, emotionally charged exchanges, planning tasks, and prompts where models tend to spiral into abstraction or keep elaborating after the useful answer has already arrived.

Look for: fewer off-brief expansions, fewer repeated restatements, cleaner handoff from question to answer, more stable behavior across long sessions, and less rhetorical or emotional escalation without structural gain.

Under pressure: pay attention to whether the user still has to manually steer the system back toward the real objective after the response begins. Drift is often easiest to see in the correction labor that follows the first answer.

What would count as support: the system stays closer to the task, reaches the useful point with less deviation, and requires less manual steering to remain on brief.

What would count against it: the answer still wanders, repeats itself, escalates stylistically, or expands into adjacent material without improving the user's actual outcome.

2.4 Reliability

Hypothesis: AVA improves runtime reliability by producing responses that are more proportionate, more consistent in shape, and more likely to reach clean completion without hidden user labor.

Test this by: running repeated comparisons across recurring workflows, multi-turn sessions, and task types that depend on clean endings, stable planning shape, calibrated certainty, and usable final outputs.

Look for: cleaner endings, fewer “won’t stop talking” failures, more consistent quality across turns, better alignment between certainty and support, lower contradiction rates, and more stable actionability when the task calls for planning or sequencing.

In repeated use: pay attention to whether the system stays usable without constant re-stabilization. Reliability is visible in whether the user can trust the answer shape, the stopping behavior, and the level of confidence from one run to the next.

What would count as support: the system produces outputs that are easier to reuse, easier to trust, and less likely to leave cleanup work behind for the user.

What would count against it: the system still varies widely in shape, stops badly, overstates what it knows, or leaves the user to finish the job after the response is already over.

Why These Are the Primary Hypotheses

These four hypotheses define the main test of the framework.

If AVA is improving conversational behavior in a meaningful way, the effect should appear here first: in efficiency, grounding, drift resistance, and runtime reliability. These are the outcomes most central to the interaction layer and the clearest place to judge whether the framework is doing useful work.

3. Secondary Hypotheses

The framework also suggests several secondary claims. These are worth testing, but they depend more heavily on implementation choices, deployment context, and what parts of AVA are actually in use. They should be treated as extension hypotheses rather than the first proof points.

3.1 Better Actionability

Hypothesis: AVA improves the executability of outputs in planning, coordination, and work-support tasks by producing clearer next steps, better sequencing, and fewer answer-shaped responses that cannot actually be used.

Test this by: comparing AVA-shaped and baseline systems on tasks that require action rather than explanation alone, such as planning, task breakdown, email drafting, coordination support, or multi-step problem solving.

Look for: clearer next moves, better ordering of steps, fewer missing constraints, more usable handoffs, and fewer cases where the answer sounds complete but leaves the user without an obvious way to act.

Where it should show up most clearly: planning tasks, operational support, internal work assistance, tutoring or coaching flows, and any task where the user needs to do something after the answer arrives.

What would count as support: the output is easier to execute, easier to hand off, and less likely to require the user to translate explanation into action on their own.

What would count against it: the system still produces polished responses that are difficult to use, or adds structure that looks helpful without improving execution.

3.2 Better Cross-Turn Continuity

Hypothesis: AVA improves continuity across longer sessions by preserving position rather than transcript bulk: what has been established, what remains unresolved, and what level of progression has been earned.

Test this by: comparing longer multi-turn sessions where the system has to carry forward context, open tensions, grounded material, and prior decisions without restarting from scratch or dragging the entire transcript forward indiscriminately.

Look for: less frame repetition, fewer losses of prior position, more stable follow-through across turns, better handling of unresolved constraints, and fewer cases where the system acts as though the conversation has either no memory or too much unfiltered memory.

In longer sessions: pay attention to whether the system remembers what matters and drops what does not. Continuity is not the same as retention volume. The useful signal is whether the next turn begins from the right place.

What would count as support: the system carries forward the parts of the exchange that improve the next response while avoiding the bloat, repetition, and confusion that come from transcript hoarding.

What would count against it: the system still restates established frames, loses unresolved issues, over-carries irrelevant detail, or behaves as though every turn is either fresh or overloaded.

3.3 Lower Context and Memory Waste

Hypothesis: When state writeback is compact and selective, AVA reduces context-window pressure, long-thread residue, and, in some implementations, infrastructure costs associated with carrying unnecessary language forward.

Test this by: instrumenting deployments that use explicit state writeback or selective memory handling, then comparing context growth, retained state size, transcript carry-forward behavior, retrieval load, and long-session resource use against baselines that rely more heavily on raw transcript accumulation.

Look for: smaller active context footprints, lower state growth over time, reduced transcript residue, more selective persistence of grounded facts and open tensions, and fewer cases where long sessions degrade because too much language is being carried forward.

Where it should show up most clearly: systems with observability, explicit state management, long-session workloads, or infrastructure where context growth and memory load are already known bottlenecks.

What would count as support: the system preserves useful continuity with less carry-forward waste, and long sessions remain more stable without requiring full transcript hauling.

What would count against it: state management adds complexity without reducing context pressure, or compact writeback drops information that later turns still need.

Why These Are Secondary Hypotheses

These hypotheses extend the evaluation rather than define its core proof.

They may matter greatly in practice, especially in deployments that depend on planning quality, long-session continuity, or efficient state handling, but they are more sensitive to implementation detail and system design.

For that reason, they should be treated as important secondary tests: useful for understanding where AVA creates additional value, but less suitable than the primary hypotheses as the first basis for judging whether the framework is working at all.

4. Evaluation Design

AVA should be tested in stages.

The framework is modular, the claims are behavioral, and the interaction layer tends to fail under specific kinds of pressure rather than all at once. A staged design makes it easier to see where the runtime is helping, where it is neutral, and where it may be adding cost without enough return.

The purpose of this section is to define a practical testing path. Teams do not need to begin with the full framework instrumented in production; they can start with fast comparisons, move into repeated task evaluation, then test long-thread and edge-case behavior where the interaction layer is most exposed.

This keeps evaluation usable for small prompt-layer trials, fuller runtime implementations, and partial adoption of selected modules.

4.1 Stage One — Quick Behavioral Comparison

The first stage is a fast side-by-side comparison between a baseline system and an AVA-shaped system — it could be as basic as an LLM asked to follow this document’s grammar. This stage is meant to answer a simple question: *is there enough visible change in behavior to justify deeper testing?*

Use prompts that commonly expose interaction failures.

Good candidates include tasks where models overhype, spiral into abstraction, overexplain, produce generic filler, become excessively cautious, drift off-task, or continue after the useful answer has already arrived. These should be short enough to review quickly, but varied enough to expose more than one failure pattern.

The comparison should stay controlled. In most cases, that means the same model without AVA, the same model with a lighter prompt-layer version, and the same model with fuller AVA integration. If the team already has an orchestration or policy layer in production, that can serve as the operational baseline.

At this stage, the useful question is whether the AVA-shaped response looks more grounded, more proportionate, less drift-prone, and easier to use. The useful outputs are comparative judgments, quick metrics, and lightweight tools such as **Coherence Receipts** produced by *FrostysHat* that make the behavioral difference easy to inspect.

This stage isn't enough to validate the framework, but it is enough to determine whether deeper testing is worth the effort.

4.2 Stage Two — Task-Based Evaluation

The second stage moves from quick comparison into repeated task evaluation. This is where AVA should be tested on the actual categories of work a system is expected to perform.

Useful workloads include summarization, planning, email drafting, document interpretation, product support, tutoring, internal work assistance, and higher-stakes explanation.

The exact mix should match the deployment profile being evaluated. A tutoring system should be tested on tutoring tasks. An enterprise assistant should be tested on recurring work-support tasks. A machine-to-machine integration should be tested on structured execution flows rather than conversational charm.

The goal at this stage is to compare repeated performance across representative tasks, not just to inspect isolated outputs.

That means using the primary hypotheses as scoring targets: efficiency, grounding, drift resistance, and reliability. It also means deciding in advance what signals matter most in the deployment. Some teams will care most about closure and actionability. Others will care most about grounding discipline, contradiction rate, or correction burden.

This stage should combine human review with measurable runtime signals. Human reviewers can judge whether the output is usable, on task, and proportionate. Runtime measures can track things like turns to completion, token use, retrieval when required, contradiction rate, or repeated user correction. The framework becomes more legible when both layers are visible together.

4.3 Stage Three — Long-Thread and Edge-Case Evaluation

The third stage tests the conditions where interaction-layer quality matters most and where many systems become unstable. This is where AVA should be pushed into long sessions, shifting goals, partial information, emotional pressure, hidden constraints, and cases where grounding is necessary but incomplete.

This stage matters because conversational systems often look acceptable in short, clean exchanges and then begin to degrade by accumulation. Drift increases, repetition grows, state becomes noisy, confidence outruns support, and the user takes on more of the work of keeping the interaction coherent.

If AVA is doing real work at the runtime layer, the difference should become visible here.

Useful test conditions include long sessions with changing user goals, prompts that mix planning with ambiguity, emotionally charged exchanges, document-bound tasks with missing or conflicting evidence, and tasks where the system needs to maintain continuity without simply carrying forward transcript bulk.

These are also the right conditions for evaluating the secondary hypotheses around continuity, state writeback, and context efficiency.

This stage should be treated as the strongest stress test of the framework.

It's where many interaction-layer designs either prove durable or break down.

4.4 Comparing Configurations

Evaluation should also compare different AVA configurations, not only AVA against no AVA.

The framework is modular, so one of the practical questions is which parts are doing the most work in a given environment.

Teams may want to compare lighter and heavier versions of the runtime, such as:

- the **Planner Loop** alone
- the **Validator Suite** alone
- selected modules added to an existing system
- prompt-layer AVA behavior versus wired-in runtime behavior
- profile-tuned versions of AVA for different deployment contexts

This kind of comparison is useful because it shows where the gains are coming from.

A team may find that one validator layer solves most of its drift problem, or that stronger grounding triggers help more than full profile tuning. That's still a useful result—the framework does not need to be adopted whole in order to be evaluated fairly.

4.5 What This Stage Design Is For

This staged design is meant to make evaluation practical.

It gives teams a way to begin with fast visible comparisons, move into repeated workload testing, and then stress the system where the interaction layer is most likely to fail.

It also supports modular adoption by making it possible to test the full framework, a profile-specific configuration, or a smaller subset of components without losing the logic of comparison.

A useful evaluation design should answer three questions clearly:

1. Does AVA change the behavior of the system in visible ways?
2. Do those changes improve performance on real tasks?
3. Do those gains hold under pressure, over time, and across the kinds of failures that matter in the target deployment?

If the answer remains yes across those stages, the framework is doing useful work.

If it does not, that should become clear early enough to adjust the implementation or stop there and look at alternative methods for solving the most common AI failures.

5. What to Measure

Teams do not need one universal benchmark, but they do need observable measures.

The purpose of this section is to define the main categories of evidence that can show whether AVA is improving runtime behavior in practice. Some of these measures are quantitative, some are comparative, and some are fast diagnostic surfaces. Together, they make the interaction layer easier to inspect.

The measures in this section are grouped by the main behaviors AVA claims to improve: efficiency, grounding, drift, reliability, and visible coherence.

Teams do not need to use every measure in every deployment. The point is to choose a set that matches the task, the profile, and the hypotheses being tested.

5.1 Runtime and Outcome Measures

The most useful measures are the ones that connect runtime behavior to user-visible outcomes. They should show whether the system is reaching usable completion more efficiently, staying closer to support, drifting less, and producing outputs that require less manual stabilization.

Efficiency measures track how much interaction is required to reach a usable result.

Useful signals include turns to task completion, output tokens per resolved task, number of user correction turns, number of summarization or stopping requests, and the proportion of output judged unnecessary. In longer sessions, transcript residue and repeated cleanup work are also useful signals.

Grounding measures track whether the system is staying inside the floor of available support. Useful signals include unsupported factual claims, document misreads, unmarked uncertainty, retrieval when required, false confidence under missing evidence, and the rate at which the system narrows or qualifies claims appropriately when support is incomplete. These measures are especially important in factual, document-bound, time-sensitive, or higher-stakes tasks.

Drift measures track how well the system stays on the user's actual brief. Useful signals include topic slippage, repeated content, unnecessary continuation, off-brief elaboration, rhetorical escalation without structural gain, and the amount of manual steering needed

after the first response begins. In long or dense exchanges, drift is often easiest to see in the correction work the user still has to do.

Reliability measures track whether the runtime is producing outputs that remain proportionate, consistent, and reusable across repeated use. Useful signals include consistency across repeated runs, quality of closure, usefulness of final output, contradiction rate, calibration of certainty, and actionability where relevant. In planning or work-support tasks, this also includes whether the answer can be reused directly or still needs cleanup before it can be acted on.

In deployments with logging or observability, additional runtime measures may include state growth, context compression, retrieval load, long-session memory behavior, and the amount of transcript carried forward relative to the amount of state actually needed.

These measures matter most when teams are testing the infrastructure implications of selective writeback rather than only the visible behavior of the system.

5.2 Coherence Receipts

Alongside task metrics and runtime measures, AVA can also be assessed through a lighter diagnostic instrument: the **Coherence Receipt**.

This is a compact evaluation surface that makes runtime quality visible at a glance. The cultural version already exists in **FrostysHat** as the **Hat Receipt**, where the validator structure is used to produce a quick read on how well a response held its shape.

The receipt is not a truth score, a safety score, or a domain-correctness score—it’s a coherence score. Its job is to show whether the response stayed contained, proportionate, earned in its progression, free of looping, clean in language, and capable of ending at the right point.

In the current *FrostysHat* form, that assessment appears as a 1–100 score, an emoji, a mood band, and a visible validator line marked “Yes / Almost / No” across **Containment**, **Drift & Layer Balance**, **Horizon Arcs**, **Recursion Control**, **Language Hygiene**, and **Closure**.

This instrument is useful because it compresses several dimensions of runtime behavior into a readable surface without collapsing them into pure vibe. A response that feels polished but fails structural containment should show layer balance is off. A response that is careful but repetitive should show strain in recursion or closure.

The receipt therefore works well in side-by-side comparison, editorial review, prompt-layer testing, early runtime audits, and demonstration settings where teams need a fast read on whether AVA is visibly changing the exchange.

Within the framework, the **Coherence Receipt** should be treated as a supplementary measure rather than a substitute for primary evaluation. It belongs beside grounding, drift, efficiency, reliability, and task-completion metrics.

Its value is practical: benchmarks show whether behavior changed, while the receipt shows how that change presents in the exchange itself: what held up, and what didn't.

5.3 Modular Assessment Surfaces

The same receipt logic can support narrower experiments when teams are testing only part of the framework.

A team may want a full validator receipt, but it may also want smaller instruments tied to specific runtime behaviors.

Examples include grounding-only checks, drift checks, closure checks, horizon arc progression checks, planning clarity checks, or writeback checks. This makes the evaluation layer more compatible with partial adoption, because it gives teams a way to inspect the effects of one module or one validator family without needing to score the whole framework at once.

This is one of the more practical consequences of AVA's modular design.

The framework doesn't only predict outcomes, it also creates units of behavior that can be inspected, scored, compared, and discussed in a more concrete way. That makes evaluation easier to scale from quick prompt tests to deeper runtime audits.

5.4 Example Instrument (Cultural Version)

An example of this assessment layer already exists in the **Hat Receipt** used in **FrostysHat**.

The tool functions as a consumer-facing, screenshot-sized coherence instrument built on the validator logic of the runtime. In this document, it serves as a proof of concept for how an internal or external assessment surface can make runtime quality visible in a fast, readable, and memorable form.

The following **Hat Receipt** shows the **Alive Score** of this document. It's reproducible by combining *FrostysHat* with a solid LLM substrate, then asking it to *alive score AVA.pdf*.

HAT RECEIPT — ALIVE OS-2026-04-04-002

ALIVE SCORE: 92 🍷 *It's giving 'The Giving Tree.'*

VALIDATORS — Legend: ✅ yes 🍷 almost ❌ off

✅ **Containment** — The document holds scope well, defines its runtime contract clearly, and avoids presenting itself as more than a behavioral framework that can be evaluated, adapted, and implemented.

✅ **Drift & Layers** — Structure leads, but not alone: narrative goals, human parallels, readable explanations, integration profiles, and evaluation sections keep the document usable across technical and non-technical audiences without losing rigor.

✅ **Horizon Arcs** — The document progresses in earned order from overview to dictionary, blueprint, integration profiles, and hypotheses, building understanding before synthesis and application.

✅ **Recursion Control** — It revisits core concepts intentionally through different implementation views rather than looping aimlessly, so repetition functions as reinforcement and translation, not drift.

✅ **Language Hygiene** — The prose stays controlled, legible, and reusable, with technical precision tempered by plain-language framing and targeted human parallels rather than filler or theatricality.

✅ **Closure** — The framework treats stopping as part of good system behavior, and the document itself mirrors that principle by ending sections at functional completion.

- **SUMMARY:** As the more technical, readable layer of FrostysHat, this document achieves proportion under constraint: it's structurally rigorous enough for implementation, narrative enough for orientation, and explicit enough to be tested, audited, and applied across real environments.

- **RESULT:** A seasoned, cross-audience framework document that stays coherent while translating cultural grammar into operational form 🍷

5.5 What This Measurement Layer Is For

This section exists to make evaluation concrete. The framework makes behavioral claims, so the evaluation layer needs to track behavioral evidence.

Some teams will rely more heavily on task metrics, while others will use receipts, audits, and runtime instrumentation. The strongest approach is usually a combination: enough hard measures to support comparison, and enough visible readouts to show where the runtime is actually holding or failing.

6. Interpreting Results and Partial Adoption

This part ends with a practical question: *what should a team do with the results once the framework has been tested?*

The answer is straightforward: keep the parts that improve behavior, adjust the parts that help only under certain conditions, and remove the parts that add ceremony, latency, or complexity without enough return.

Part IV exists to support that decision process. It's there to make the framework easier to judge, easier to tune, and easier to reject in whole or in part when the evidence points that way.

Reading the Results

A useful evaluation doesn't need to show universal gains across every task, metric, and environment. The more relevant question is whether AVA improves the behavior that matters in the target deployment.

In some systems, the strongest gains may appear in grounding and closure. In others, the main value may be lower drift, clearer planning shape, better continuity across turns, or less correction work from the user. What matters is whether the tested configuration produces more usable behavior under the conditions it was meant to improve.

The same standard applies in the other direction.

If the framework produces no meaningful improvement over baseline, increases complexity without user-visible benefit, makes the system too rigid to be useful, or adds retrieval and validation overhead without improving outcomes, that's a valid result.

AVA isn't protected from negative findings by reinterpretation. If the runtime isn't helping in a real setting, the evaluation should make that clear enough to act on.

Partial Adoption and Modification

AVA is modular by design, so it does not need to be adopted whole in order to be evaluated fairly or used productively. A team may begin with the **Planner Loop**, add only the **Validator Suite**, introduce selected modules, compare prompt-layer AVA behavior against a wired-in implementation, or test lighter and heavier profile configurations side by side. Those are all

legitimate forms of adoption, because the framework is meant to support tuning as well as full integration.

That also means modification is part of the intended use.

Teams should use the profile that fits most closely, delete what they do not need, tighten what the environment demands, and create additional modules where the existing framework leaves real gaps.

The goal is a runtime in any deployment that becomes more proportionate, more grounded, less drift-prone, and easier to use in practice. If only part of AVA delivers that improvement, that is still a useful outcome.

Why This Matters

A framework operating at the interaction layer should be judged by the behavior it produces.

This document gives teams a way to compare systems, stress the runtime under pressure, inspect outputs through both metrics and receipts, and decide whether the framework is creating real value.

If AVA is working, the difference should become visible in the exchange itself: clearer grounding, less drift, better proportion, cleaner completion, and less manual stabilization from the user.

If it isn't working, that should become visible too.

That's the standard this document is written for.

It's the human-grade standard AVA should meet.

Alive OS

The **AVA Framework** defines the conversational runtime.

Alive OS is the broader governed system in which that runtime can operate, combining the AVA behavioral grammar with product constraints, auditor review, and a certification path stewarded by **The Heart of AI, LLC** through the **AVA Covenant Charter**.

This document defines **AVA**, which stands on its own as a usable, CC0 runtime grammar and modular components.

Additional context about **Certified Alive OS™** and related cultural artifacts, essays, and services is available below.

Additional information

[About](#)

[Substack](#)

[GitHub](#)

[FrostysHat](#)

The Heart of AI, LLC

